

Question 1

Minimization of Negative Log Likelihood with Softmax outputs

We are given the negative log likelihood function that we want to minimize:

$$J(w) = -\log(\prod_n \prod_{m=0}^9 p(t_n = m|x_n; w))$$

where there are 0 to 9 classes.

Now let us define that $p(t_m|x; w)$ is a categorical probability distribution that is :

$$p(t_m|x; w) = y_m(x; w)^{t_m}$$

following the constraints that $t_m \in \{0, 1\}$ and $\sum_m t_m = 1$

As we know that we can interpret the output of the softmax as the probabilities that a certain set of features belongs to a certain class. Let us use the probability nodes as softmax output nodes

$$y_i(x; w) = \frac{\exp^{y_i(x; w)}}{\sum_{j=0}^m \exp^{y_j(x; w)}}$$

This ensures that sum of probability of each feature sample belonging to one of the 10 classes equals to 1

The negative log likelihood can be reduced to :

$$J(w) = -\sum_n \sum_{m=0}^9 \log p(t_n = m|x_n; w)$$

$$J(w) = -\sum_n \sum_{m=0}^9 t_m^{(n)} \log(y_m^{(n)}(x^{(n)}; w))$$

Now the nature of curve of negative log likelihood is :




Figure: *The loss function reaches infinity when input is 0, and reaches 0 when input is 1.*

This means for the curve to reach its minimum we have to generate the likelihood of the correct label as close to 1 as possible.

Now noting that t_m is one-hot encoded vector thus reducing all terms of $J(w)$ to 0 other than the target(true) class say 'k' for every feature vector i.e.

$$J(w) = - \sum_n 1 * \log(y_k^{(n)}(x^{(n)}; w)) = - \log(\prod_n (y_k^{(n)}(x^{(n)}; w)))$$

Thus the weights must be trained in such a way that the likelihood softmax(probability) of the target label is maximum for each feature vector (i.e. close to 1).

Hence it is proved that a neural network to maximize the log likelihood of label is one that has softmax output nodes and minimizes the criterion function of the negative log probability of training data set

Interpretation of L2 regularization

Let us assume that our input is linearly correlated with our output using the following equation

$$y_i = wx_i + \epsilon$$

where w is the set of weights and ϵ is the random noise that also follows normal distribution with 0 mean and constant variance σ^2 .

This gives us the likelihood that we want to maximize:

$$\prod_{i=1}^N \mathcal{N}(y_i; x_i, w, \sigma^2)$$

Now, let us assume that w has the Gaussian prior $\mathcal{N}(0, \alpha^{-1})$

From Bayes Rule we know:

$$P(w|D) = \frac{P(D|w)P(w)}{P(D)}$$

$$P(w|D) \approx P(D|w)P(w)$$

Thus the posterior probability is:

$$P(w|D) = \prod_{i=1}^N \mathcal{N}(y_i; x_i, w, \sigma^2) \mathcal{N}(0, \alpha^{-1})$$

Taking logarithm of the posterior:

$$\log P(w|D) = \sum_{i=1}^N \log \mathcal{N}(y_i; x_i, w, \sigma^2) + \log \mathcal{N}(0, \alpha^{-1}) + \text{constant}$$

Ignoring the constant we need to maximize over w

$$-\frac{1}{\sigma^2} \sum_{i=1}^N (y_i - wx_i)^2 - \alpha w^2$$

This is the same as minimizing the cost function with L2 regularization with α being the regularization constant.

Thus we have proved optimizing model weights to minimize loss function with L2 regularization is equivalent to finding the weights that are most likely under a posterior distribution evaluated using Bayes rule, with a zero-mean independent Gaussian weights prior

```
In [1]: import tensorflow as tf
import numpy as np
```

```
In [2]: from keras.datasets import mnist

(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

Using TensorFlow backend.

```
In [3]: x_train.shape,y_train.shape
```

```
Out[3]: ((60000, 28, 28), (60000,))
```

```
In [4]: x_test.shape,y_test.shape
```

```
Out[4]: ((10000, 28, 28), (10000,))
```

Reshaping and normalizing the data

```
In [5]: x_train=x_train.reshape((60000,784))
```

```
x_test=x_test.reshape((10000,784))
```

```
In [6]: from sklearn.preprocessing import StandardScaler
```

```
sc=StandardScaler()
```

```
x_train=sc.fit_transform(x_train)
```

```
x_test=sc.fit_transform(x_test)
```

Generating 1000 images with 100 images per class

```
In [8]: def generate_minidata(x,y):
```

```
    mini=[]
```

```
    for i in range(10):
```

```
        count=0
```

```
        for j in range(len(x)):
```

```
            if y[j]==i:
```

```
                mini.append(np.array([i]+list(x[j])))
```

```
                count+=1
```

```
            if count==100:
```

```
                break
```

```
    mini=np.array(mini)
```

```
    np.random.shuffle(mini)
```

```
    #return(mini[:,0].reshape((len(mini),1)),mini[:,1:])
```

```
    return(mini[:,0],mini[:,1:])
```

```
In [9]: mini_trainlabels,mini_train=generate_minidata(x_train,y_train)
```

```
In [10]: mini_testlabels,mini_test=generate_minidata(x_test,y_test)
```

```
In [11]: mini_train.shape,mini_trainlabels.shape
```

```
Out[11]: ((1000, 784), (1000,))
```

```
In [12]: mini_test.shape,mini_testlabels.shape
```

```
Out[12]: ((1000, 784), (1000,))
```

One hot encoding

```
In [15]: from sklearn.preprocessing import OneHotEncoder
         ohe = OneHotEncoder()
         train_labels = ohe.fit_transform(mini_trainlabels.reshape((1000,1))
         ).toarray()
         test_labels = ohe.fit_transform(mini_testlabels.reshape((1000,1))).
         toarray()
```

```
In [16]: print(mini_test.shape,test_labels.shape)

(1000, 784) (1000, 10)
```

```
In [17]: print(mini_train.shape,train_labels.shape)

(1000, 784) (1000, 10)
```

Question 2)a) 1 Hidden Layer

```
In [31]: def neural_net_1hidden(l2):

    model = tf.keras.Sequential()

    if l2==True:
        model.add(tf.keras.layers.Dense(30, input_dim=784,activation='sigmoid',kernel_regularizer=tf.keras.regularizers.l2(5)))

    else:
        model.add(tf.keras.layers.Dense(30, input_dim=784,activation='sigmoid'))

    model.add(tf.keras.layers.Dense(10, activation='softmax'))

    model.summary()

    adam=tf.keras.optimizers.Adam(learning_rate=0.1)
    model.compile(loss='categorical_crossentropy', optimizer=adam,
metrics=[ 'accuracy' ])

    return model
```

One hidden layer without regularization

```
In [19]: model=neural_net_1hidden(False)
history = model.fit(mini_train, train_labels, epochs=30, batch_size=10,validation_data=(mini_test,test_labels))
```

Model: "sequential"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| dense (Dense) | (None, 30) | 23550 |
| dense_1 (Dense) | (None, 10) | 310 |

Total params: 23,860

Trainable params: 23,860

Non-trainable params: 0

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

1000/1000 [=====] - 1s 1ms/sample - loss: 1.3298 - accuracy: 0.5870 - val_loss: 1.2903 - val_accuracy: 0.5730

Epoch 2/30

1000/1000 [=====] - 0s 316us/sample - loss: 0.8941 - accuracy: 0.7260 - val_loss: 1.1346 - val_accuracy: 0.6490

Epoch 3/30

1000/1000 [=====] - 0s 317us/sample - loss: 0.8941 - accuracy: 0.7260 - val_loss: 1.1346 - val_accuracy: 0.6490

```
s: 0.8239 - accuracy: 0.7640 - val_loss: 1.1899 - val_accuracy: 0.6460
Epoch 4/30
1000/1000 [=====] - 0s 312us/sample - loss: 0.7272 - accuracy: 0.7710 - val_loss: 1.1487 - val_accuracy: 0.6830
Epoch 5/30
1000/1000 [=====] - 0s 309us/sample - loss: 0.8568 - accuracy: 0.7420 - val_loss: 1.0889 - val_accuracy: 0.6760
Epoch 6/30
1000/1000 [=====] - 0s 327us/sample - loss: 0.8361 - accuracy: 0.7670 - val_loss: 1.2460 - val_accuracy: 0.6770
Epoch 7/30
1000/1000 [=====] - 0s 310us/sample - loss: 0.8100 - accuracy: 0.7710 - val_loss: 1.1949 - val_accuracy: 0.6740
Epoch 8/30
1000/1000 [=====] - 0s 301us/sample - loss: 0.7046 - accuracy: 0.8110 - val_loss: 1.1376 - val_accuracy: 0.6910
Epoch 9/30
1000/1000 [=====] - 0s 375us/sample - loss: 0.7463 - accuracy: 0.7980 - val_loss: 1.3716 - val_accuracy: 0.6540
Epoch 10/30
1000/1000 [=====] - 0s 328us/sample - loss: 0.6351 - accuracy: 0.8250 - val_loss: 1.3194 - val_accuracy: 0.6680
Epoch 11/30
1000/1000 [=====] - 0s 308us/sample - loss: 0.6398 - accuracy: 0.8300 - val_loss: 1.1581 - val_accuracy: 0.6810
Epoch 12/30
1000/1000 [=====] - 0s 313us/sample - loss: 0.7264 - accuracy: 0.8040 - val_loss: 1.2404 - val_accuracy: 0.6870
Epoch 13/30
1000/1000 [=====] - 0s 308us/sample - loss: 0.6077 - accuracy: 0.8230 - val_loss: 1.3915 - val_accuracy: 0.6760
Epoch 14/30
1000/1000 [=====] - 0s 310us/sample - loss: 0.5773 - accuracy: 0.8400 - val_loss: 1.9983 - val_accuracy: 0.6000
Epoch 15/30
1000/1000 [=====] - 0s 306us/sample - loss: 0.6211 - accuracy: 0.8330 - val_loss: 1.3181 - val_accuracy: 0.7010
Epoch 16/30
1000/1000 [=====] - 0s 306us/sample - loss: 0.5186 - accuracy: 0.8460 - val_loss: 1.4768 - val_accuracy: 0.
```

```
6900
Epoch 17/30
1000/1000 [=====] - 0s 308us/sample - loss: 0.6365 - accuracy: 0.8260 - val_loss: 1.2859 - val_accuracy: 0.7180
Epoch 18/30
1000/1000 [=====] - 0s 308us/sample - loss: 0.6268 - accuracy: 0.8310 - val_loss: 1.2983 - val_accuracy: 0.7160
Epoch 19/30
1000/1000 [=====] - 0s 301us/sample - loss: 0.5686 - accuracy: 0.8380 - val_loss: 1.4851 - val_accuracy: 0.6900
Epoch 20/30
1000/1000 [=====] - 0s 330us/sample - loss: 0.5521 - accuracy: 0.8500 - val_loss: 1.2269 - val_accuracy: 0.7240
Epoch 21/30
1000/1000 [=====] - 0s 301us/sample - loss: 0.5034 - accuracy: 0.8640 - val_loss: 1.3665 - val_accuracy: 0.7180
Epoch 22/30
1000/1000 [=====] - 0s 303us/sample - loss: 0.5890 - accuracy: 0.8380 - val_loss: 1.4347 - val_accuracy: 0.6960
Epoch 23/30
1000/1000 [=====] - 0s 337us/sample - loss: 0.5691 - accuracy: 0.8460 - val_loss: 1.4923 - val_accuracy: 0.6800
Epoch 24/30
1000/1000 [=====] - 0s 321us/sample - loss: 0.5675 - accuracy: 0.8580 - val_loss: 1.5891 - val_accuracy: 0.6740
Epoch 25/30
1000/1000 [=====] - 0s 417us/sample - loss: 0.5334 - accuracy: 0.8490 - val_loss: 1.4392 - val_accuracy: 0.6800
Epoch 26/30
1000/1000 [=====] - 0s 431us/sample - loss: 0.6023 - accuracy: 0.8260 - val_loss: 1.8166 - val_accuracy: 0.6560
Epoch 27/30
1000/1000 [=====] - 0s 324us/sample - loss: 0.5137 - accuracy: 0.8680 - val_loss: 1.6423 - val_accuracy: 0.7040
Epoch 28/30
1000/1000 [=====] - 0s 333us/sample - loss: 0.5182 - accuracy: 0.8610 - val_loss: 1.5842 - val_accuracy: 0.7010
Epoch 29/30
1000/1000 [=====] - 0s 318us/sample - loss: 0.5649 - accuracy: 0.8660 - val_loss: 1.4517 - val_accuracy: 0.7210
```





```
Epoch 30/30  
1000/1000 [=====] - 0s 359us/sample - los  
s: 0.5091 - accuracy: 0.8650 - val_loss: 1.6785 - val_accuracy: 0.  
6880
```

```
In [20]: import matplotlib.pyplot as plt  
  
def model_training_plot():  
  
    # Plot training & validation loss values  
    plt.plot(history.history['loss'])  
    plt.plot(history.history['val_loss'])  
    plt.title('Model loss')  
    plt.ylabel('Loss')  
    plt.xlabel('Epoch')  
    plt.legend(['Train', 'Test'], loc='upper left')  
    plt.show()  
  
    # Plot training & validation accuracy values  
    plt.plot(history.history['accuracy'])  
    plt.plot(history.history['val_accuracy'])  
    plt.title('Model accuracy')  
    plt.ylabel('Accuracy')  
    plt.xlabel('Epoch')  
    plt.legend(['Train', 'Test'], loc='upper left')  
    plt.show()
```

Plot of criterion categorical cross entropy loss function

```
In [21]: model_training_plot()
```



```
In [22]: def for_error_plot(model):

    train_errors=[]
    test_errors=[]
    train_accuracy=[]
    test_accuracy=[]

    for i in range(30):

        model.fit(mini_train, train_labels, epochs=1, batch_size=10
        )

        predictions_train=model.predict_classes(mini_train)

        predictions_test=model.predict_classes(mini_test)

        te=0
        tr=0
        for i in range(1000):
            if train_labels[i][predictions_train[i]]==1:
                tr+=1
            if test_labels[i][predictions_test[i]]==1:
                te+=1

        train_errors.append(1-tr/1000)
        test_errors.append(1-te/1000)

        train_accuracy.append(tr/1000)
        test_accuracy.append(te/1000)

    return train_errors,test_errors,train_accuracy,test_accuracy
```

```
In [23]: model=neural_net_1hidden(False)
         traine,teste,traina,testa=for_error_plot(model)
```

Model: "sequential_1"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| dense_2 (Dense) | (None, 30) | 23550 |
| dense_3 (Dense) | (None, 10) | 310 |

Total params: 23,860

Trainable params: 23,860

Non-trainable params: 0

Train on 1000 samples

1000/1000 [=====] - 1s 525us/sample - los
s: 1.3550 - accuracy: 0.6010

```
Train on 1000 samples
1000/1000 [=====] - 0s 190us/sample - loss: 0.8118 - accuracy: 0.7450
Train on 1000 samples
1000/1000 [=====] - 0s 188us/sample - loss: 0.8000 - accuracy: 0.7720
Train on 1000 samples
1000/1000 [=====] - 0s 209us/sample - loss: 0.7834 - accuracy: 0.7830
Train on 1000 samples
1000/1000 [=====] - 0s 204us/sample - loss: 0.6142 - accuracy: 0.8240
Train on 1000 samples
1000/1000 [=====] - 0s 236us/sample - loss: 0.6590 - accuracy: 0.8080
Train on 1000 samples
1000/1000 [=====] - 0s 213us/sample - loss: 0.5990 - accuracy: 0.8320
Train on 1000 samples
1000/1000 [=====] - 0s 188us/sample - loss: 0.6900 - accuracy: 0.8000
Train on 1000 samples
1000/1000 [=====] - 0s 194us/sample - loss: 0.6767 - accuracy: 0.8140
Train on 1000 samples
1000/1000 [=====] - 0s 193us/sample - loss: 0.6917 - accuracy: 0.8270
Train on 1000 samples
1000/1000 [=====] - 0s 183us/sample - loss: 0.6282 - accuracy: 0.8360
Train on 1000 samples
1000/1000 [=====] - 0s 180us/sample - loss: 0.5449 - accuracy: 0.8400
Train on 1000 samples
1000/1000 [=====] - 0s 177us/sample - loss: 0.5866 - accuracy: 0.8440
Train on 1000 samples
1000/1000 [=====] - 0s 181us/sample - loss: 0.5944 - accuracy: 0.8400
Train on 1000 samples
1000/1000 [=====] - 0s 184us/sample - loss: 0.6462 - accuracy: 0.8310
Train on 1000 samples
1000/1000 [=====] - 0s 183us/sample - loss: 0.7162 - accuracy: 0.8270
Train on 1000 samples
1000/1000 [=====] - 0s 185us/sample - loss: 0.7007 - accuracy: 0.8120
Train on 1000 samples
1000/1000 [=====] - 0s 180us/sample - loss: 0.6469 - accuracy: 0.8470
Train on 1000 samples
1000/1000 [=====] - 0s 182us/sample - loss:
```


```
s: 0.6453 - accuracy: 0.8280
Train on 1000 samples
1000/1000 [=====] - 0s 181us/sample - loss: 0.6324 - accuracy: 0.8440
Train on 1000 samples
1000/1000 [=====] - 0s 205us/sample - loss: 0.6496 - accuracy: 0.8350
Train on 1000 samples
1000/1000 [=====] - 0s 243us/sample - loss: 0.6711 - accuracy: 0.8280
Train on 1000 samples
1000/1000 [=====] - 0s 210us/sample - loss: 0.5601 - accuracy: 0.8580
Train on 1000 samples
1000/1000 [=====] - 0s 178us/sample - loss: 0.5569 - accuracy: 0.8570
Train on 1000 samples
1000/1000 [=====] - 0s 194us/sample - loss: 0.5587 - accuracy: 0.8500
Train on 1000 samples
1000/1000 [=====] - 0s 185us/sample - loss: 0.5770 - accuracy: 0.8680
Train on 1000 samples
1000/1000 [=====] - 0s 179us/sample - loss: 0.5083 - accuracy: 0.8690
Train on 1000 samples
1000/1000 [=====] - 0s 179us/sample - loss: 0.5445 - accuracy: 0.8620
Train on 1000 samples
1000/1000 [=====] - 0s 213us/sample - loss: 0.5288 - accuracy: 0.8670
Train on 1000 samples
1000/1000 [=====] - 0s 213us/sample - loss: 0.4683 - accuracy: 0.8690
```

```
In [24]: import matplotlib.pyplot as plt
def zero_one_error_plot(x,y,a,b):
    # Plot training & validation loss values
    plt.plot(x)
    plt.plot(y)
    plt.title('Model Error')
    plt.ylabel('Error')
    plt.xlabel('Epoch')
    plt.legend(['Train', 'Test'], loc='upper left')
    plt.show()

    plt.plot(a)
    plt.plot(b)
    plt.title('Model Accuracy')
    plt.ylabel('Accuracy')
    plt.xlabel('Epoch')
    plt.legend(['Train', 'Test'], loc='upper left')
    plt.show()
```

Plot of zero-one error

```
In [25]: zero_one_error_plot(traine, teste, traina, testa)
```



```
In [26]: def learning_rate_plot(params,model):

    learning_rate=[]

    for i in range(30):
        old_weights=model.get_weights()

        model.fit(mini_train, train_labels, epochs=1, batch_size=10
        )

        new_weights=model.get_weights()

        summation=0

        for j in range(len(new_weights)):
            diff=np.sum(np.absolute((old_weights[j]-new_weights[j])
            /new_weights[j]))
            summation+=diff

        learning_rate.append(summation/params)

    # Plot learning rate
    plt.plot(learning_rate)
    plt.title('Learning rate')
    plt.ylabel('Rate')
    plt.xlabel('Epoch')
    plt.show()
```

Learning speed of hidden layer

```
In [28]: model=neural_net_1hidden(False)
        learning_rate_plot(23860,model)
```

Model: "sequential_2"

| Layer (type) | Output Shape | Param # |
|-----------------|--------------|---------|
| dense_4 (Dense) | (None, 30) | 23550 |
| dense_5 (Dense) | (None, 10) | 310 |

Total params: 23,860

Trainable params: 23,860

Non-trainable params: 0

Train on 1000 samples

1000/1000 [=====] - 1s 525us/sample - loss: 1.3328 - accuracy: 0.5860

Train on 1000 samples


1000/1000 [=====] - 0s 194us/sample - loss: 0.8527 - accuracy: 0.7260


```
Train on 1000 samples
1000/1000 [=====] - 0s 180us/sample - loss: 0.7924 - accuracy: 0.7580
Train on 1000 samples
1000/1000 [=====] - 0s 178us/sample - loss: 0.7943 - accuracy: 0.7660
Train on 1000 samples
1000/1000 [=====] - 0s 181us/sample - loss: 0.7105 - accuracy: 0.8000
Train on 1000 samples
1000/1000 [=====] - 0s 177us/sample - loss: 0.6744 - accuracy: 0.8120
Train on 1000 samples
1000/1000 [=====] - 0s 181us/sample - loss: 0.7136 - accuracy: 0.7940
Train on 1000 samples
1000/1000 [=====] - 0s 179us/sample - loss: 0.5927 - accuracy: 0.8290
Train on 1000 samples
1000/1000 [=====] - 0s 179us/sample - loss: 0.6010 - accuracy: 0.8250
Train on 1000 samples
1000/1000 [=====] - 0s 182us/sample - loss: 0.5567 - accuracy: 0.8310
Train on 1000 samples
1000/1000 [=====] - 0s 182us/sample - loss: 0.6206 - accuracy: 0.8300
Train on 1000 samples
1000/1000 [=====] - 0s 183us/sample - loss: 0.6232 - accuracy: 0.8190
Train on 1000 samples
1000/1000 [=====] - 0s 184us/sample - loss: 0.6614 - accuracy: 0.8210
Train on 1000 samples
1000/1000 [=====] - 0s 180us/sample - loss: 0.5912 - accuracy: 0.8470
Train on 1000 samples
1000/1000 [=====] - 0s 191us/sample - loss: 0.6023 - accuracy: 0.8310
Train on 1000 samples
1000/1000 [=====] - 0s 181us/sample - loss: 0.5788 - accuracy: 0.8460
Train on 1000 samples
1000/1000 [=====] - 0s 181us/sample - loss: 0.4843 - accuracy: 0.8600
Train on 1000 samples
1000/1000 [=====] - 0s 190us/sample - loss: 0.5590 - accuracy: 0.8540
Train on 1000 samples
1000/1000 [=====] - 0s 181us/sample - loss: 0.5239 - accuracy: 0.8590
Train on 1000 samples
1000/1000 [=====] - 0s 185us/sample - loss:
```

```

s: 0.6009 - accuracy: 0.8370
Train on 1000 samples
1000/1000 [=====] - 0s 184us/sample - los
s: 0.5834 - accuracy: 0.8400
Train on 1000 samples
1000/1000 [=====] - 0s 180us/sample - los
s: 0.5561 - accuracy: 0.8540
Train on 1000 samples
1000/1000 [=====] - 0s 181us/sample - los
s: 0.5381 - accuracy: 0.8490
Train on 1000 samples
1000/1000 [=====] - 0s 180us/sample - los
s: 0.5531 - accuracy: 0.8320
Train on 1000 samples
1000/1000 [=====] - 0s 181us/sample - los
s: 0.5594 - accuracy: 0.8340
Train on 1000 samples
1000/1000 [=====] - 0s 180us/sample - los
s: 0.4907 - accuracy: 0.8630
Train on 1000 samples
1000/1000 [=====] - 0s 180us/sample - los
s: 0.5608 - accuracy: 0.8510
Train on 1000 samples
1000/1000 [=====] - 0s 178us/sample - los
s: 0.5077 - accuracy: 0.8570
Train on 1000 samples
1000/1000 [=====] - 0s 180us/sample - los
s: 0.4772 - accuracy: 0.8670
Train on 1000 samples
1000/1000 [=====] - 0s 181us/sample - los
s: 0.4610 - accuracy: 0.8750

```



One hidden layer with regularization

With criterion function loss

```
In [32]: model=neural_net_1hidden(True)
history = model.fit(mini_train, train_labels, epochs=30, batch_size
=10,validation_data=(mini_test,test_labels))
model_training_plot()
```

Model: "sequential_5"

| Layer (type) | Output Shape | Param # |
|--------------------------|--------------|---------|
| dense_9 (Dense) | (None, 30) | 23550 |
| dense_10 (Dense) | (None, 10) | 310 |
| Total params: 23,860 | | |
| Trainable params: 23,860 | | |
| Non-trainable params: 0 | | |

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

1000/1000 [=====] - 1s 818us/sample - loss: 30.7975 - accuracy: 0.1440 - val_loss: 3.2870 - val_accuracy: 0.2530

Epoch 2/30

1000/1000 [=====] - 0s 335us/sample - loss: 4.1948 - accuracy: 0.1780 - val_loss: 4.4483 - val_accuracy: 0.0910

Epoch 3/30

1000/1000 [=====] - 0s 419us/sample - loss: 4.5490 - accuracy: 0.1870 - val_loss: 6.2060 - val_accuracy: 0.1130

Epoch 4/30

1000/1000 [=====] - 0s 358us/sample - loss: 5.4089 - accuracy: 0.1870 - val_loss: 5.4600 - val_accuracy: 0.1670

Epoch 5/30

1000/1000 [=====] - 0s 317us/sample - loss: 5.4986 - accuracy: 0.2120 - val_loss: 6.5071 - val_accuracy: 0.1910

Epoch 6/30

1000/1000 [=====] - 0s 328us/sample - loss: 6.0101 - accuracy: 0.2060 - val_loss: 6.7398 - val_accuracy: 0.1310

Epoch 7/30

1000/1000 [=====] - 0s 387us/sample - loss: 6.3535 - accuracy: 0.2180 - val_loss: 5.8183 - val_accuracy: 0.1170

Epoch 8/30


1000/1000 [=====] - 0s 352us/sample - loss: 5.4076 - accuracy: 0.1910 - val_loss: 6.0906 - val_accuracy: 0.1410


```
Epoch 9/30
1000/1000 [=====] - 0s 321us/sample - loss: 6.2441 - accuracy: 0.2140 - val_loss: 8.2925 - val_accuracy: 0.2000
Epoch 10/30
1000/1000 [=====] - 0s 326us/sample - loss: 6.1735 - accuracy: 0.1930 - val_loss: 4.6071 - val_accuracy: 0.2560
Epoch 11/30
1000/1000 [=====] - 0s 414us/sample - loss: 5.7798 - accuracy: 0.2010 - val_loss: 6.4946 - val_accuracy: 0.1270
Epoch 12/30
1000/1000 [=====] - 0s 314us/sample - loss: 6.8121 - accuracy: 0.2010 - val_loss: 6.9213 - val_accuracy: 0.1250
Epoch 13/30
1000/1000 [=====] - 0s 370us/sample - loss: 6.4988 - accuracy: 0.2040 - val_loss: 6.9691 - val_accuracy: 0.1220
Epoch 14/30
1000/1000 [=====] - 0s 355us/sample - loss: 7.1023 - accuracy: 0.2110 - val_loss: 8.6227 - val_accuracy: 0.1580
Epoch 15/30
1000/1000 [=====] - 0s 378us/sample - loss: 6.6713 - accuracy: 0.2150 - val_loss: 6.2033 - val_accuracy: 0.1370
Epoch 16/30
1000/1000 [=====] - 0s 329us/sample - loss: 5.3896 - accuracy: 0.1550 - val_loss: 4.9290 - val_accuracy: 0.1740
Epoch 17/30
1000/1000 [=====] - 0s 328us/sample - loss: 5.6952 - accuracy: 0.2010 - val_loss: 4.9898 - val_accuracy: 0.1510
Epoch 18/30
1000/1000 [=====] - 0s 348us/sample - loss: 5.3386 - accuracy: 0.1740 - val_loss: 4.5789 - val_accuracy: 0.1820
Epoch 19/30
1000/1000 [=====] - 0s 335us/sample - loss: 5.8689 - accuracy: 0.2000 - val_loss: 6.0154 - val_accuracy: 0.1500
Epoch 20/30
1000/1000 [=====] - 0s 370us/sample - loss: 6.8612 - accuracy: 0.2010 - val_loss: 7.9184 - val_accuracy: 0.1450
Epoch 21/30
1000/1000 [=====] - 0s 358us/sample - loss: 6.2756 - accuracy: 0.1740 - val_loss: 4.9567 - val_accuracy: 0.2140
Epoch 22/30
```

```

1000/1000 [=====] - 0s 347us/sample - loss: 5.9909 - accuracy: 0.1940 - val_loss: 4.7433 - val_accuracy: 0.2170
Epoch 23/30
1000/1000 [=====] - 0s 313us/sample - loss: 6.4548 - accuracy: 0.1800 - val_loss: 5.8111 - val_accuracy: 0.0940
Epoch 24/30
1000/1000 [=====] - 0s 425us/sample - loss: 6.2695 - accuracy: 0.1920 - val_loss: 5.7078 - val_accuracy: 0.1020
Epoch 25/30
1000/1000 [=====] - 0s 387us/sample - loss: 5.4674 - accuracy: 0.1850 - val_loss: 5.0711 - val_accuracy: 0.1470
Epoch 26/30
1000/1000 [=====] - 0s 408us/sample - loss: 6.1370 - accuracy: 0.1690 - val_loss: 7.3002 - val_accuracy: 0.1190
Epoch 27/30
1000/1000 [=====] - 0s 367us/sample - loss: 6.0303 - accuracy: 0.1790 - val_loss: 5.1454 - val_accuracy: 0.1750
Epoch 28/30
1000/1000 [=====] - 0s 313us/sample - loss: 6.3046 - accuracy: 0.2140 - val_loss: 4.6457 - val_accuracy: 0.3300
Epoch 29/30
1000/1000 [=====] - 0s 331us/sample - loss: 6.7093 - accuracy: 0.2050 - val_loss: 6.0724 - val_accuracy: 0.1810
Epoch 30/30
1000/1000 [=====] - 0s 317us/sample - loss: 6.1461 - accuracy: 0.2180 - val_loss: 6.0274 - val_accuracy: 0.1430

```





Zero-one Error

```
In [33]: model=neural_net_1hidden(True)
         traine,teste,traina,testa=for_error_plot(model)
         zero_one_error_plot(traine,teste,traina,testa)
```

Model: "sequential_6"

| Layer (type) | Output Shape | Param # |
|------------------|--------------|---------|
| dense_11 (Dense) | (None, 30) | 23550 |
| dense_12 (Dense) | (None, 10) | 310 |

=====
 Total params: 23,860
 Trainable params: 23,860
 Non-trainable params: 0


```
Train on 1000 samples
1000/1000 [=====] - 1s 756us/sample - loss: 30.8283 - accuracy: 0.1350
Train on 1000 samples
1000/1000 [=====] - 0s 259us/sample - loss: 4.3015 - accuracy: 0.1640
Train on 1000 samples
1000/1000 [=====] - 0s 238us/sample - loss: 5.2170 - accuracy: 0.2010
Train on 1000 samples
1000/1000 [=====] - 0s 186us/sample - loss: 4.8569 - accuracy: 0.1790
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 5.3937 - accuracy: 0.2070
Train on 1000 samples
```


```
1000/1000 [=====] - 0s 185us/sample - loss: 5.8270 - accuracy: 0.1980
Train on 1000 samples
1000/1000 [=====] - 0s 183us/sample - loss: 5.7690 - accuracy: 0.2310
Train on 1000 samples
1000/1000 [=====] - 0s 188us/sample - loss: 6.7065 - accuracy: 0.2510
Train on 1000 samples
1000/1000 [=====] - 0s 187us/sample - loss: 6.5117 - accuracy: 0.2100
Train on 1000 samples
1000/1000 [=====] - 0s 197us/sample - loss: 5.9477 - accuracy: 0.2310
Train on 1000 samples
1000/1000 [=====] - 0s 219us/sample - loss: 7.1525 - accuracy: 0.2470
Train on 1000 samples
1000/1000 [=====] - 0s 189us/sample - loss: 7.4980 - accuracy: 0.2010
Train on 1000 samples
1000/1000 [=====] - 0s 199us/sample - loss: 5.6581 - accuracy: 0.2220
Train on 1000 samples
1000/1000 [=====] - 0s 204us/sample - loss: 5.8502 - accuracy: 0.1720
Train on 1000 samples
1000/1000 [=====] - 0s 189us/sample - loss: 6.3494 - accuracy: 0.2100
Train on 1000 samples
1000/1000 [=====] - 0s 237us/sample - loss: 6.4638 - accuracy: 0.1990
Train on 1000 samples
1000/1000 [=====] - 0s 222us/sample - loss: 5.8281 - accuracy: 0.2160
Train on 1000 samples
1000/1000 [=====] - 0s 242us/sample - loss: 6.0723 - accuracy: 0.1920
Train on 1000 samples
1000/1000 [=====] - 0s 251us/sample - loss: 5.7845 - accuracy: 0.2230
Train on 1000 samples
1000/1000 [=====] - 0s 244us/sample - loss: 6.9948 - accuracy: 0.2010
Train on 1000 samples
1000/1000 [=====] - 0s 231us/sample - loss: 6.9080 - accuracy: 0.1870
Train on 1000 samples
1000/1000 [=====] - 0s 220us/sample - loss: 5.6190 - accuracy: 0.2100
Train on 1000 samples
1000/1000 [=====] - 0s 206us/sample - loss: 5.6348 - accuracy: 0.1820
```

```

Train on 1000 samples
1000/1000 [=====] - 0s 197us/sample - loss: 5.9712 - accuracy: 0.1790
Train on 1000 samples
1000/1000 [=====] - 0s 235us/sample - loss: 6.0057 - accuracy: 0.1990
Train on 1000 samples
1000/1000 [=====] - ETA: 0s - loss: 5.8145 - accuracy: 0.21 - 0s 210us/sample - loss: 5.8027 - accuracy: 0.2150
Train on 1000 samples
1000/1000 [=====] - 0s 221us/sample - loss: 6.6944 - accuracy: 0.1860
Train on 1000 samples
1000/1000 [=====] - 0s 186us/sample - loss: 5.9219 - accuracy: 0.2030
Train on 1000 samples
1000/1000 [=====] - 0s 247us/sample - loss: 5.9162 - accuracy: 0.2150
Train on 1000 samples
1000/1000 [=====] - 0s 220us/sample - loss: 5.8164 - accuracy: 0.2100

```





Learning rate plot

```
In [34]: model=neural_net_1hidden(True)
         learning_rate_plot(23860,model)
```

Model: "sequential_7"

| Layer (type) | Output Shape | Param # |
|------------------|--------------|---------|
| dense_13 (Dense) | (None, 30) | 23550 |
| dense_14 (Dense) | (None, 10) | 310 |

Total params: 23,860

Trainable params: 23,860

Non-trainable params: 0

Train on 1000 samples

1000/1000 [=====] - 1s 572us/sample - loss: 30.8666 - accuracy: 0.1610

Train on 1000 samples

1000/1000 [=====] - 0s 246us/sample - loss: 4.2672 - accuracy: 0.1750

Train on 1000 samples

1000/1000 [=====] - 0s 184us/sample - loss: 5.1633 - accuracy: 0.2090

Train on 1000 samples

1000/1000 [=====] - 0s 185us/sample - loss: 5.4903 - accuracy: 0.2090

Train on 1000 samples

1000/1000 [=====] - 0s 193us/sample - loss: 5.6340 - accuracy: 0.1900

Train on 1000 samples


1000/1000 [=====] - 0s 204us/sample - loss: 5.6340 - accuracy: 0.1900

```
s: 5.4936 - accuracy: 0.2080
Train on 1000 samples
1000/1000 [=====] - 0s 183us/sample - los
s: 7.0238 - accuracy: 0.2390
Train on 1000 samples
1000/1000 [=====] - 0s 185us/sample - los
s: 6.9018 - accuracy: 0.2250
Train on 1000 samples
1000/1000 [=====] - 0s 185us/sample - los
s: 6.2076 - accuracy: 0.2170
Train on 1000 samples
1000/1000 [=====] - 0s 185us/sample - los
s: 6.7087 - accuracy: 0.2140
Train on 1000 samples
1000/1000 [=====] - 0s 184us/sample - los
s: 5.8844 - accuracy: 0.2360
Train on 1000 samples
1000/1000 [=====] - 0s 186us/sample - los
s: 5.8784 - accuracy: 0.2300
Train on 1000 samples
1000/1000 [=====] - 0s 245us/sample - los
s: 6.5338 - accuracy: 0.2330
Train on 1000 samples
1000/1000 [=====] - 0s 226us/sample - los
s: 6.4696 - accuracy: 0.2190
Train on 1000 samples
1000/1000 [=====] - 0s 226us/sample - los
s: 6.2606 - accuracy: 0.2100
Train on 1000 samples
1000/1000 [=====] - 0s 206us/sample - los
s: 7.1551 - accuracy: 0.2230
Train on 1000 samples
1000/1000 [=====] - 0s 238us/sample - los
s: 6.6390 - accuracy: 0.1950
Train on 1000 samples
1000/1000 [=====] - 0s 239us/sample - los
s: 7.7188 - accuracy: 0.2360
Train on 1000 samples
1000/1000 [=====] - 0s 230us/sample - los
s: 6.6714 - accuracy: 0.2460
Train on 1000 samples
1000/1000 [=====] - 0s 256us/sample - los
s: 6.9349 - accuracy: 0.2280
Train on 1000 samples
1000/1000 [=====] - 0s 226us/sample - los
s: 6.7569 - accuracy: 0.1880
Train on 1000 samples
1000/1000 [=====] - 0s 220us/sample - los
s: 7.3080 - accuracy: 0.2300
Train on 1000 samples
1000/1000 [=====] - 0s 236us/sample - los
s: 6.8204 - accuracy: 0.2140
Train on 1000 samples
```

```

1000/1000 [=====] - 0s 219us/sample - loss: 6.7142 - accuracy: 0.2160
Train on 1000 samples
1000/1000 [=====] - 0s 242us/sample - loss: 7.0781 - accuracy: 0.2330
Train on 1000 samples
1000/1000 [=====] - 0s 213us/sample - loss: 6.5930 - accuracy: 0.2090
Train on 1000 samples
1000/1000 [=====] - 0s 191us/sample - loss: 6.4572 - accuracy: 0.2260
Train on 1000 samples
1000/1000 [=====] - 0s 182us/sample - loss: 6.8373 - accuracy: 0.2180
Train on 1000 samples
1000/1000 [=====] - 0s 197us/sample - loss: 6.1812 - accuracy: 0.1930
Train on 1000 samples
1000/1000 [=====] - 0s 199us/sample - loss: 6.0630 - accuracy: 0.1970

```



Question 2)b) 2 and 3 hidden Layers

```
In [35]: def neural_net_2hidden(l2):

    model = tf.keras.Sequential()
    if l2==True:
        model.add(tf.keras.layers.Dense(30, input_dim=784,activation
n='sigmoid',kernel_regularizer=tf.keras.regularizers.l2(5)))
        model.add(tf.keras.layers.Dense(30, input_dim=30,activation
='sigmoid',kernel_regularizer=tf.keras.regularizers.l2(5)))
    else:
        model.add(tf.keras.layers.Dense(30, input_dim=784,activation
n='sigmoid'))
        model.add(tf.keras.layers.Dense(30, input_dim=30,activation
='sigmoid'))

    model.add(tf.keras.layers.Dense(10, activation='softmax'))

    model.summary()

    adam=tf.keras.optimizers.Adam(learning_rate=0.1)
    model.compile(loss='categorical_crossentropy', optimizer=adam,
metrics=[ 'accuracy' ])

    return model
```

Two hidden layers without regularization

```
In [36]: model=neural_net_2hidden(False)
history = model.fit(mini_train, train_labels, epochs=30, batch_size
=10,validation_data=(mini_test,test_labels))
```

Model: "sequential_8"

| Layer (type) | Output Shape | Param # |
|--------------------------|--------------|---------|
| dense_15 (Dense) | (None, 30) | 23550 |
| dense_16 (Dense) | (None, 30) | 930 |
| dense_17 (Dense) | (None, 10) | 310 |
| Total params: 24,790 | | |
| Trainable params: 24,790 | | |
| Non-trainable params: 0 | | |

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

1000/1000 [=====] - 1s 826us/sample - loss: 1.6937 - accuracy: 0.4290 - val_loss: 1.3850 - val_accuracy: 0.5760

Epoch 2/30

1000/1000 [=====] - 0s 328us/sample - loss

```
s: 1.1888 - accuracy: 0.5890 - val_loss: 1.2996 - val_accuracy: 0.5610
Epoch 3/30
1000/1000 [=====] - 0s 316us/sample - loss: 0.9741 - accuracy: 0.6900 - val_loss: 1.2834 - val_accuracy: 0.5980
Epoch 4/30
1000/1000 [=====] - 0s 319us/sample - loss: 0.9653 - accuracy: 0.6890 - val_loss: 1.2727 - val_accuracy: 0.5980
Epoch 5/30
1000/1000 [=====] - 0s 373us/sample - loss: 0.9281 - accuracy: 0.7160 - val_loss: 1.4058 - val_accuracy: 0.5930
Epoch 6/30
1000/1000 [=====] - 0s 428us/sample - loss: 0.9608 - accuracy: 0.7140 - val_loss: 1.2359 - val_accuracy: 0.6080
Epoch 7/30
1000/1000 [=====] - 0s 348us/sample - loss: 0.8343 - accuracy: 0.7340 - val_loss: 1.1676 - val_accuracy: 0.6580
Epoch 8/30
1000/1000 [=====] - 0s 321us/sample - loss: 0.8357 - accuracy: 0.7480 - val_loss: 1.3857 - val_accuracy: 0.5990
Epoch 9/30
1000/1000 [=====] - 0s 338us/sample - loss: 0.7733 - accuracy: 0.7560 - val_loss: 1.2879 - val_accuracy: 0.6490
Epoch 10/30
1000/1000 [=====] - 0s 414us/sample - loss: 0.7853 - accuracy: 0.7730 - val_loss: 1.3084 - val_accuracy: 0.6160
Epoch 11/30
1000/1000 [=====] - 0s 332us/sample - loss: 0.8449 - accuracy: 0.7470 - val_loss: 1.3589 - val_accuracy: 0.6190
Epoch 12/30
1000/1000 [=====] - 0s 334us/sample - loss: 0.8092 - accuracy: 0.7500 - val_loss: 1.2941 - val_accuracy: 0.6540
Epoch 13/30
1000/1000 [=====] - 0s 324us/sample - loss: 0.8074 - accuracy: 0.7440 - val_loss: 1.1574 - val_accuracy: 0.6660
Epoch 14/30
1000/1000 [=====] - 0s 322us/sample - loss: 0.7708 - accuracy: 0.7590 - val_loss: 1.4231 - val_accuracy: 0.6060
Epoch 15/30
1000/1000 [=====] - 0s 416us/sample - loss: 0.8277 - accuracy: 0.7440 - val_loss: 1.2309 - val_accuracy: 0.
```

```
6680
Epoch 16/30
1000/1000 [=====] - 0s 320us/sample - loss: 0.7890 - accuracy: 0.7710 - val_loss: 1.1239 - val_accuracy: 0.6630
Epoch 17/30
1000/1000 [=====] - 0s 337us/sample - loss: 0.7663 - accuracy: 0.7610 - val_loss: 1.2652 - val_accuracy: 0.6470
Epoch 18/30
1000/1000 [=====] - 0s 381us/sample - loss: 0.7632 - accuracy: 0.7670 - val_loss: 1.1659 - val_accuracy: 0.6650
Epoch 19/30
1000/1000 [=====] - 0s 360us/sample - loss: 0.7140 - accuracy: 0.7640 - val_loss: 1.1964 - val_accuracy: 0.6860
Epoch 20/30
1000/1000 [=====] - 0s 372us/sample - loss: 0.7335 - accuracy: 0.7700 - val_loss: 1.3293 - val_accuracy: 0.6170
Epoch 21/30
1000/1000 [=====] - 0s 421us/sample - loss: 0.7198 - accuracy: 0.7770 - val_loss: 1.2356 - val_accuracy: 0.6600
Epoch 22/30
1000/1000 [=====] - 0s 398us/sample - loss: 0.6204 - accuracy: 0.7960 - val_loss: 1.3508 - val_accuracy: 0.6710
Epoch 23/30
1000/1000 [=====] - 0s 413us/sample - loss: 0.5583 - accuracy: 0.8190 - val_loss: 1.2347 - val_accuracy: 0.6650
Epoch 24/30
1000/1000 [=====] - 0s 376us/sample - loss: 0.6145 - accuracy: 0.8000 - val_loss: 1.2766 - val_accuracy: 0.6790
Epoch 25/30
1000/1000 [=====] - 0s 318us/sample - loss: 0.6433 - accuracy: 0.7920 - val_loss: 1.2652 - val_accuracy: 0.6430
Epoch 26/30
1000/1000 [=====] - 0s 345us/sample - loss: 0.6458 - accuracy: 0.7980 - val_loss: 1.2674 - val_accuracy: 0.7030
Epoch 27/30
1000/1000 [=====] - 0s 414us/sample - loss: 0.6359 - accuracy: 0.8170 - val_loss: 1.4330 - val_accuracy: 0.6350
Epoch 28/30
1000/1000 [=====] - 0s 383us/sample - loss: 0.6125 - accuracy: 0.8110 - val_loss: 1.1482 - val_accuracy: 0.6780
```

Epoch 29/30


1000/1000 [=====] - 0s 331us/sample - loss: 0.6119 - accuracy: 0.8100 - val_loss: 1.3694 - val_accuracy: 0.6550

Epoch 30/30

1000/1000 [=====] - 0s 325us/sample - loss: 0.6618 - accuracy: 0.8020 - val_loss: 1.3414 - val_accuracy: 0.6530

Plot for criterion function loss

```
In [37]: model_training_plot()
```



Plot for zero-one error

```
In [38]: model=neural_net_2hidden(False)
         traine,teste,traina,testa=for_error_plot(model)
         zero_one_error_plot(traine,teste,traina,testa)
```

Model: "sequential_9"

| Layer (type) | Output Shape | Param # |
|------------------|--------------|---------|
| dense_18 (Dense) | (None, 30) | 23550 |
| dense_19 (Dense) | (None, 30) | 930 |
| dense_20 (Dense) | (None, 10) | 310 |

=====
Total params: 24,790

Trainable params: 24,790

Non-trainable params: 0

=====
Train on 1000 samples

1000/1000 [=====] - 1s 558us/sample - loss: 1.8175 - accuracy: 0.3470

Train on 1000 samples

1000/1000 [=====] - 0s 188us/sample - loss: 1.1615 - accuracy: 0.5880

Train on 1000 samples

1000/1000 [=====] - 0s 194us/sample - loss: 0.9782 - accuracy: 0.6580

Train on 1000 samples

1000/1000 [=====] - 0s 193us/sample - loss: 0.9635 - accuracy: 0.6710

Train on 1000 samples

1000/1000 [=====] - 0s 188us/sample - loss: 0.9301 - accuracy: 0.7080

Train on 1000 samples

1000/1000 [=====] - 0s 193us/sample - loss: 0.9206 - accuracy: 0.7150

Train on 1000 samples

1000/1000 [=====] - 0s 195us/sample - loss: 0.8934 - accuracy: 0.7140

Train on 1000 samples

1000/1000 [=====] - 0s 225us/sample - loss: 0.8474 - accuracy: 0.7390

Train on 1000 samples

1000/1000 [=====] - 0s 194us/sample - loss: 0.7925 - accuracy: 0.7540

Train on 1000 samples

1000/1000 [=====] - 0s 195us/sample - loss: 0.7751 - accuracy: 0.7620

Train on 1000 samples

1000/1000 [=====] - 0s 192us/sample - loss: 0.7839 - accuracy: 0.7720

Train on 1000 samples

1000/1000 [=====] - 0s 196us/sample - loss: 0.7839 - accuracy: 0.7720

Train on 1000 samples

1000/1000 [=====] - 0s 196us/sample - loss: 0.7839 - accuracy: 0.7720

Train on 1000 samples

1000/1000 [=====] - 0s 196us/sample - loss: 0.7839 - accuracy: 0.7720

Train on 1000 samples

1000/1000 [=====] - 0s 196us/sample - loss: 0.7839 - accuracy: 0.7720

Train on 1000 samples

1000/1000 [=====] - 0s 196us/sample - loss: 0.7839 - accuracy: 0.7720

Train on 1000 samples

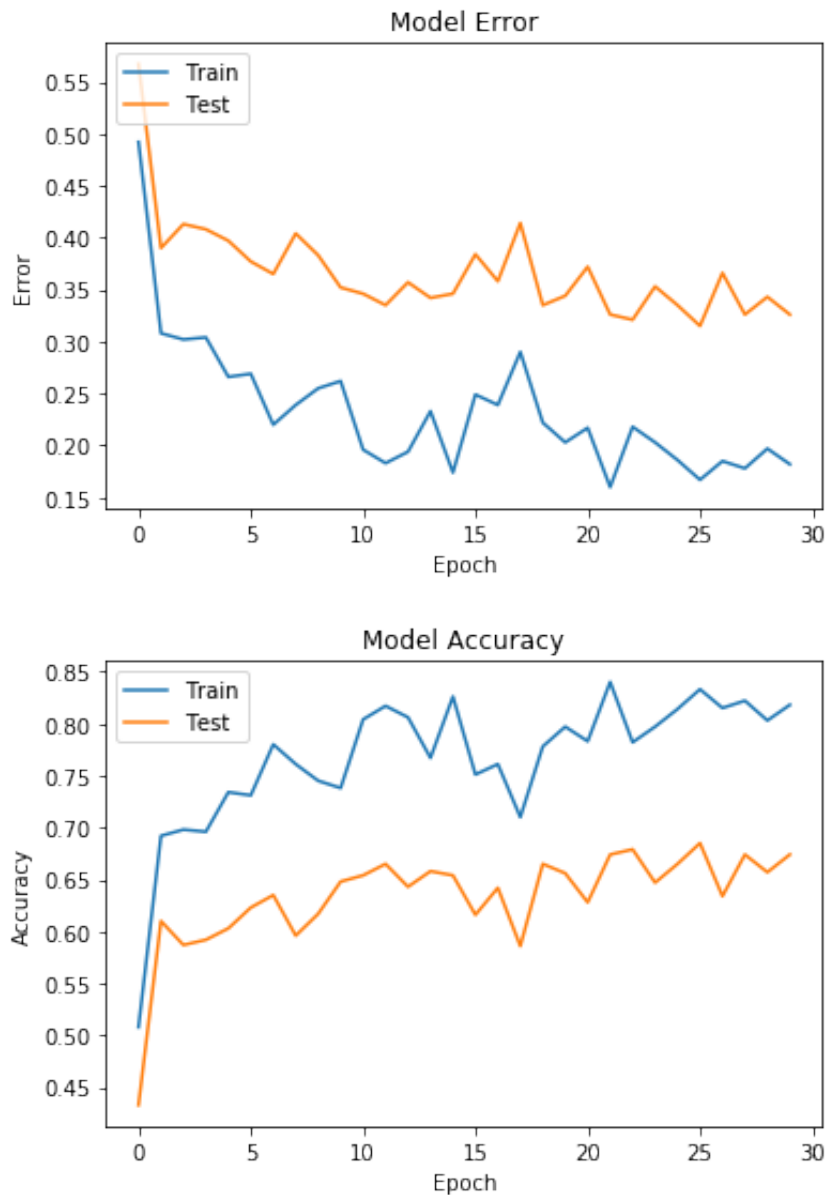
1000/1000 [=====] - 0s 196us/sample - loss: 0.7839 - accuracy: 0.7720

Train on 1000 samples

1000/1000 [=====] - 0s 196us/sample - loss: 0.7839 - accuracy: 0.7720


```
s: 0.7849 - accuracy: 0.7540
Train on 1000 samples
1000/1000 [=====] - 0s 193us/sample - los
s: 0.7600 - accuracy: 0.7760
Train on 1000 samples
1000/1000 [=====] - 0s 195us/sample - los
s: 0.7600 - accuracy: 0.7510
Train on 1000 samples
1000/1000 [=====] - 0s 191us/sample - los
s: 0.7074 - accuracy: 0.7770
Train on 1000 samples
1000/1000 [=====] - 0s 193us/sample - los
s: 0.7434 - accuracy: 0.7780
Train on 1000 samples
1000/1000 [=====] - 0s 189us/sample - los
s: 0.7563 - accuracy: 0.7660
Train on 1000 samples
1000/1000 [=====] - 0s 258us/sample - los
s: 0.8655 - accuracy: 0.7580
Train on 1000 samples
1000/1000 [=====] - 0s 229us/sample - los
s: 0.8545 - accuracy: 0.7510
Train on 1000 samples
1000/1000 [=====] - 0s 192us/sample - los
s: 0.7556 - accuracy: 0.7520
Train on 1000 samples
1000/1000 [=====] - 0s 197us/sample - los
s: 0.7085 - accuracy: 0.7850
Train on 1000 samples
1000/1000 [=====] - 0s 196us/sample - los
s: 0.6538 - accuracy: 0.7960
Train on 1000 samples
1000/1000 [=====] - 0s 198us/sample - los
s: 0.6791 - accuracy: 0.7920
Train on 1000 samples
1000/1000 [=====] - 0s 193us/sample - los
s: 0.6676 - accuracy: 0.8030
Train on 1000 samples
1000/1000 [=====] - 0s 195us/sample - los
s: 0.6636 - accuracy: 0.7970
Train on 1000 samples
1000/1000 [=====] - 0s 191us/sample - los
s: 0.6071 - accuracy: 0.8100
Train on 1000 samples
1000/1000 [=====] - 0s 192us/sample - los
s: 0.5883 - accuracy: 0.8070
Train on 1000 samples
1000/1000 [=====] - 0s 252us/sample - los
s: 0.6070 - accuracy: 0.8140
Train on 1000 samples
1000/1000 [=====] - 0s 240us/sample - los
s: 0.6154 - accuracy: 0.8040
Train on 1000 samples
```

1000/1000 [=====] - 0s 196us/sample - loss: 0.6613 - accuracy: 0.7890



Plot for learning rate

```
In [40]: model=neural_net_2hidden(False)
         learning_rate_plot(24790,model)
```


Model: "sequential_11"

| Layer (type) | Output Shape | Param # |
|------------------|--------------|---------|
| dense_24 (Dense) | (None, 30) | 23550 |
| dense_25 (Dense) | (None, 30) | 930 |
| dense_26 (Dense) | (None, 10) | 310 |

Total params: 24,790
Trainable params: 24,790
Non-trainable params: 0

Train on 1000 samples
1000/1000 [=====] - 1s 638us/sample - loss: 1.5793 - accuracy: 0.4430
Train on 1000 samples
1000/1000 [=====] - 0s 208us/sample - loss: 1.0365 - accuracy: 0.6630
Train on 1000 samples
1000/1000 [=====] - 0s 229us/sample - loss: 0.9198 - accuracy: 0.6940
Train on 1000 samples
1000/1000 [=====] - 0s 238us/sample - loss: 0.8485 - accuracy: 0.7250
Train on 1000 samples
1000/1000 [=====] - 0s 212us/sample - loss: 0.8280 - accuracy: 0.7160
Train on 1000 samples
1000/1000 [=====] - 0s 194us/sample - loss: 0.7396 - accuracy: 0.7490
Train on 1000 samples
1000/1000 [=====] - 0s 196us/sample - loss: 0.7668 - accuracy: 0.7180
Train on 1000 samples
1000/1000 [=====] - 0s 206us/sample - loss: 0.8473 - accuracy: 0.7290
Train on 1000 samples
1000/1000 [=====] - 0s 188us/sample - loss: 0.7289 - accuracy: 0.7600
Train on 1000 samples
1000/1000 [=====] - 0s 188us/sample - loss: 0.7743 - accuracy: 0.7510
Train on 1000 samples
1000/1000 [=====] - 0s 189us/sample - loss: 0.8649 - accuracy: 0.7320
Train on 1000 samples
1000/1000 [=====] - 0s 188us/sample - loss: 0.7400 - accuracy: 0.7660
Train on 1000 samples
1000/1000 [=====] - 0s 191us/sample - loss: 0.7700 - accuracy: 0.7470
Train on 1000 samples
1000/1000 [=====] - 0s 189us/sample - loss: 0.7770 - accuracy: 0.7560
Train on 1000 samples
1000/1000 [=====] - 0s 189us/sample - loss: 0.7830 - accuracy: 0.7470
Train on 1000 samples
1000/1000 [=====] - 0s 195us/sample - loss: 0.7971 - accuracy: 0.7480
Train on 1000 samples

```
1000/1000 [=====] - 0s 187us/sample - loss: 0.7495 - accuracy: 0.7570
Train on 1000 samples
1000/1000 [=====] - 0s 189us/sample - loss: 0.7103 - accuracy: 0.7900
Train on 1000 samples
1000/1000 [=====] - 0s 201us/sample - loss: 0.7099 - accuracy: 0.7940
Train on 1000 samples
1000/1000 [=====] - 0s 190us/sample - loss: 0.6896 - accuracy: 0.7780
Train on 1000 samples
1000/1000 [=====] - 0s 193us/sample - loss: 0.6564 - accuracy: 0.7850
Train on 1000 samples
1000/1000 [=====] - 0s 190us/sample - loss: 0.7390 - accuracy: 0.7920
Train on 1000 samples
1000/1000 [=====] - 0s 186us/sample - loss: 0.6696 - accuracy: 0.8010
Train on 1000 samples
1000/1000 [=====] - 0s 189us/sample - loss: 0.6313 - accuracy: 0.8160
Train on 1000 samples
1000/1000 [=====] - 0s 193us/sample - loss: 0.6275 - accuracy: 0.8160
Train on 1000 samples
1000/1000 [=====] - 0s 247us/sample - loss: 0.6753 - accuracy: 0.7880
Train on 1000 samples
1000/1000 [=====] - 0s 216us/sample - loss: 0.7381 - accuracy: 0.7760
Train on 1000 samples
1000/1000 [=====] - 0s 250us/sample - loss: 0.6682 - accuracy: 0.7990
Train on 1000 samples
1000/1000 [=====] - 0s 223us/sample - loss: 0.6817 - accuracy: 0.7940
Train on 1000 samples
1000/1000 [=====] - 0s 248us/sample - loss: 0.5951 - accuracy: 0.8070
```



Two hidden layers with regularization

```
In [41]: model=neural_net_2hidden(True)
history = model.fit(mini_train, train_labels, epochs=30, batch_size
=10,validation_data=(mini_test,test_labels))
```

Model: "sequential_12"

| Layer (type) | Output Shape | Param # |
|------------------|--------------|---------|
| dense_27 (Dense) | (None, 30) | 23550 |
| dense_28 (Dense) | (None, 30) | 930 |
| dense_29 (Dense) | (None, 10) | 310 |

Total params: 24,790

Trainable params: 24,790

Non-trainable params: 0

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

1000/1000 [=====] - 1s 886us/sample - loss: 35.4528 - accuracy: 0.0970 - val_loss: 2.3400 - val_accuracy: 0.1000

Epoch 2/30

1000/1000 [=====] - 0s 380us/sample - loss: 2.3353 - accuracy: 0.1010 - val_loss: 2.3166 - val_accuracy: 0.1000

Epoch 3/30

1000/1000 [=====] - 0s 438us/sample - loss: 2.3390 - accuracy: 0.1040 - val_loss: 2.3211 - val_accuracy: 0.1000

Epoch 4/30


```
1000/1000 [=====] - 0s 418us/sample - loss: 2.3368 - accuracy: 0.0850 - val_loss: 2.3255 - val_accuracy: 0.1000
Epoch 5/30
1000/1000 [=====] - 0s 380us/sample - loss: 2.3283 - accuracy: 0.1210 - val_loss: 2.3461 - val_accuracy: 0.1000
Epoch 6/30
1000/1000 [=====] - 0s 348us/sample - loss: 2.3304 - accuracy: 0.0960 - val_loss: 2.3214 - val_accuracy: 0.1000
Epoch 7/30
1000/1000 [=====] - 0s 417us/sample - loss: 2.3339 - accuracy: 0.0820 - val_loss: 2.3073 - val_accuracy: 0.1000
Epoch 8/30
1000/1000 [=====] - 0s 387us/sample - loss: 2.3242 - accuracy: 0.0900 - val_loss: 2.3154 - val_accuracy: 0.1000
Epoch 9/30
1000/1000 [=====] - 0s 423us/sample - loss: 2.3284 - accuracy: 0.1040 - val_loss: 2.3150 - val_accuracy: 0.1000
Epoch 10/30
1000/1000 [=====] - 0s 348us/sample - loss: 2.3276 - accuracy: 0.1070 - val_loss: 2.3255 - val_accuracy: 0.1000
Epoch 11/30
1000/1000 [=====] - 0s 335us/sample - loss: 2.3235 - accuracy: 0.0970 - val_loss: 2.3189 - val_accuracy: 0.1000
Epoch 12/30
1000/1000 [=====] - 0s 386us/sample - loss: 2.3323 - accuracy: 0.0860 - val_loss: 2.3299 - val_accuracy: 0.1000
Epoch 13/30
1000/1000 [=====] - 0s 330us/sample - loss: 2.8774 - accuracy: 0.1000 - val_loss: 2.7886 - val_accuracy: 0.1000
Epoch 14/30
1000/1000 [=====] - 0s 410us/sample - loss: 2.8972 - accuracy: 0.1100 - val_loss: 2.8664 - val_accuracy: 0.1000
Epoch 15/30
1000/1000 [=====] - 0s 417us/sample - loss: 2.8798 - accuracy: 0.0860 - val_loss: 3.0443 - val_accuracy: 0.1000
Epoch 16/30
1000/1000 [=====] - 0s 326us/sample - loss: 3.1523 - accuracy: 0.0940 - val_loss: 3.1739 - val_accuracy: 0.1000
Epoch 17/30
1000/1000 [=====] - 0s 331us/sample - loss:
```

```
s: 3.1223 - accuracy: 0.0960 - val_loss: 3.1636 - val_accuracy: 0.1000
Epoch 18/30
1000/1000 [=====] - 0s 338us/sample - loss: 3.2697 - accuracy: 0.1000 - val_loss: 3.3040 - val_accuracy: 0.1000
Epoch 19/30
1000/1000 [=====] - 0s 324us/sample - loss: 3.2387 - accuracy: 0.0900 - val_loss: 3.1925 - val_accuracy: 0.1000
Epoch 20/30
1000/1000 [=====] - 0s 325us/sample - loss: 3.2253 - accuracy: 0.0830 - val_loss: 3.3379 - val_accuracy: 0.1000
Epoch 21/30
1000/1000 [=====] - 0s 322us/sample - loss: 3.2587 - accuracy: 0.0890 - val_loss: 3.1757 - val_accuracy: 0.1000
Epoch 22/30
1000/1000 [=====] - 0s 324us/sample - loss: 3.1113 - accuracy: 0.1150 - val_loss: 3.1259 - val_accuracy: 0.1000
Epoch 23/30
1000/1000 [=====] - 0s 364us/sample - loss: 3.2360 - accuracy: 0.0940 - val_loss: 3.2885 - val_accuracy: 0.1000
Epoch 24/30
1000/1000 [=====] - 0s 402us/sample - loss: 3.2744 - accuracy: 0.1070 - val_loss: 3.2503 - val_accuracy: 0.1000
Epoch 25/30
1000/1000 [=====] - 0s 336us/sample - loss: 3.2355 - accuracy: 0.1220 - val_loss: 3.2200 - val_accuracy: 0.1000
Epoch 26/30
1000/1000 [=====] - 0s 336us/sample - loss: 3.2211 - accuracy: 0.0950 - val_loss: 3.2609 - val_accuracy: 0.1000
Epoch 27/30
1000/1000 [=====] - 0s 334us/sample - loss: 3.1975 - accuracy: 0.1000 - val_loss: 3.2188 - val_accuracy: 0.1000
Epoch 28/30
1000/1000 [=====] - 0s 324us/sample - loss: 3.2180 - accuracy: 0.0990 - val_loss: 3.2444 - val_accuracy: 0.1000
Epoch 29/30
1000/1000 [=====] - 0s 327us/sample - loss: 3.2976 - accuracy: 0.0870 - val_loss: 3.1887 - val_accuracy: 0.1000
Epoch 30/30
1000/1000 [=====] - 0s 331us/sample - loss: 3.2200 - accuracy: 0.0830 - val_loss: 3.2762 - val_accuracy: 0.1000
```

1000

Plot for criterion function loss

```
In [42]: model_training_plot()
```

**Plot for zero-one error**

```
In [43]: model=neural_net_2hidden(True)
         traine,teste,traina,testa=for_error_plot(model)
         zero_one_error_plot(traine,teste,traina,testa)
```

Model: "sequential_13"

| Layer (type) | Output Shape | Param # |
|--------------|--------------|---------|
| ===== | | |

| | | |
|------------------|------------|-------|
| dense_30 (Dense) | (None, 30) | 23550 |
|------------------|------------|-------|

| | | |
|------------------|------------|-----|
| dense_31 (Dense) | (None, 30) | 930 |
|------------------|------------|-----|

| | | |
|------------------|------------|-----|
| dense_32 (Dense) | (None, 10) | 310 |
|------------------|------------|-----|

=====

Total params: 24,790

Trainable params: 24,790

Non-trainable params: 0

Train on 1000 samples

1000/1000 [=====] - 1s 665us/sample - loss: 35.3744 - accuracy: 0.1210

Train on 1000 samples

1000/1000 [=====] - 0s 225us/sample - loss: 2.3323 - accuracy: 0.0900

Train on 1000 samples

1000/1000 [=====] - 0s 237us/sample - loss: 2.3320 - accuracy: 0.0820

Train on 1000 samples

1000/1000 [=====] - 0s 236us/sample - loss: 2.3367 - accuracy: 0.0850

Train on 1000 samples

1000/1000 [=====] - 0s 251us/sample - loss: 2.3323 - accuracy: 0.0870

Train on 1000 samples

1000/1000 [=====] - 0s 208us/sample - loss: 2.3313 - accuracy: 0.0930

Train on 1000 samples

1000/1000 [=====] - 0s 200us/sample - loss: 2.3274 - accuracy: 0.0960

Train on 1000 samples

1000/1000 [=====] - 0s 193us/sample - loss: 2.3274 - accuracy: 0.0940

Train on 1000 samples

1000/1000 [=====] - 0s 195us/sample - loss: 2.3275 - accuracy: 0.0980

Train on 1000 samples

1000/1000 [=====] - 0s 203us/sample - loss: 2.3219 - accuracy: 0.0950

Train on 1000 samples

1000/1000 [=====] - 0s 200us/sample - loss: 2.3320 - accuracy: 0.0890

Train on 1000 samples

1000/1000 [=====] - 0s 196us/sample - loss: 2.3270 - accuracy: 0.0840

Train on 1000 samples


1000/1000 [=====] - 0s 196us/sample - loss: 2.8895 - accuracy: 0.0970

Train on 1000 samples

1000/1000 [=====] - 0s 199us/sample - loss: 2.9074 - accuracy: 0.0870

Train on 1000 samples

```
1000/1000 [=====] - 0s 195us/sample - loss: 2.8571 - accuracy: 0.0920
Train on 1000 samples
1000/1000 [=====] - 0s 199us/sample - loss: 3.1341 - accuracy: 0.1060
Train on 1000 samples
1000/1000 [=====] - 0s 203us/sample - loss: 3.1350 - accuracy: 0.0900
Train on 1000 samples
1000/1000 [=====] - 0s 193us/sample - loss: 3.2785 - accuracy: 0.1000
Train on 1000 samples
1000/1000 [=====] - 0s 194us/sample - loss: 3.2268 - accuracy: 0.1010
Train on 1000 samples
1000/1000 [=====] - 0s 194us/sample - loss: 3.2006 - accuracy: 0.0960
Train on 1000 samples
1000/1000 [=====] - 0s 197us/sample - loss: 3.2948 - accuracy: 0.0890
Train on 1000 samples
1000/1000 [=====] - 0s 197us/sample - loss: 3.1292 - accuracy: 0.0840
Train on 1000 samples
1000/1000 [=====] - 0s 194us/sample - loss: 3.2618 - accuracy: 0.0920
Train on 1000 samples
1000/1000 [=====] - 0s 196us/sample - loss: 3.2727 - accuracy: 0.1060
Train on 1000 samples
1000/1000 [=====] - 0s 195us/sample - loss: 3.2283 - accuracy: 0.0890
Train on 1000 samples
1000/1000 [=====] - 0s 196us/sample - loss: 3.2242 - accuracy: 0.0990
Train on 1000 samples
1000/1000 [=====] - 0s 209us/sample - loss: 3.1738 - accuracy: 0.0950
Train on 1000 samples
1000/1000 [=====] - 0s 279us/sample - loss: 3.2659 - accuracy: 0.0820
Train on 1000 samples
1000/1000 [=====] - 0s 227us/sample - loss: 3.2749 - accuracy: 0.0780
Train on 1000 samples
1000/1000 [=====] - 0s 198us/sample - loss: 3.2193 - accuracy: 0.0930
```



Plot for learning rate

```
In [44]: model=neural_net_2hidden(True)
         learning_rate_plot(24790,model)
```

Model: "sequential_14"

| Layer (type) | Output Shape | Param # |
|------------------|--------------|---------|
| dense_33 (Dense) | (None, 30) | 23550 |
| dense_34 (Dense) | (None, 30) | 930 |
| dense_35 (Dense) | (None, 10) | 310 |


Total params: 24,790

Trainable params: 24,790

Non-trainable params: 0

```
Train on 1000 samples
1000/1000 [=====] - 1s 637us/sample - loss: 35.4035 - accuracy: 0.1010
Train on 1000 samples
1000/1000 [=====] - 0s 195us/sample - loss: 2.3344 - accuracy: 0.0970
Train on 1000 samples
1000/1000 [=====] - 0s 195us/sample - loss: 2.3330 - accuracy: 0.1030
Train on 1000 samples
1000/1000 [=====] - 0s 198us/sample - loss: 2.3329 - accuracy: 0.0750
Train on 1000 samples
1000/1000 [=====] - 0s 237us/sample - loss: 2.3275 - accuracy: 0.1060
Train on 1000 samples
1000/1000 [=====] - 0s 215us/sample - loss: 2.3257 - accuracy: 0.1010
Train on 1000 samples
1000/1000 [=====] - 0s 237us/sample - loss: 2.3267 - accuracy: 0.0920
Train on 1000 samples
1000/1000 [=====] - 0s 238us/sample - loss: 2.3286 - accuracy: 0.0910
Train on 1000 samples
1000/1000 [=====] - 0s 190us/sample - loss: 2.3293 - accuracy: 0.1010
Train on 1000 samples
1000/1000 [=====] - 0s 189us/sample - loss: 2.3273 - accuracy: 0.1020
Train on 1000 samples
1000/1000 [=====] - 0s 201us/sample - loss: 2.3290 - accuracy: 0.0950
Train on 1000 samples
1000/1000 [=====] - 0s 188us/sample - loss: 2.3300 - accuracy: 0.1010
Train on 1000 samples
1000/1000 [=====] - 0s 186us/sample - loss: 2.8664 - accuracy: 0.1010
Train on 1000 samples
1000/1000 [=====] - 0s 187us/sample - loss: 2.8807 - accuracy: 0.0910
Train on 1000 samples
1000/1000 [=====] - 0s 187us/sample - loss: 2.8765 - accuracy: 0.0910
Train on 1000 samples
1000/1000 [=====] - 0s 185us/sample - loss: 3.1454 - accuracy: 0.0850
Train on 1000 samples
1000/1000 [=====] - 0s 187us/sample - loss: 3.1379 - accuracy: 0.0820
```

```
Train on 1000 samples
1000/1000 [=====] - 0s 187us/sample - loss: 3.2949 - accuracy: 0.0850
Train on 1000 samples
1000/1000 [=====] - 0s 205us/sample - loss: 3.1835 - accuracy: 0.1110
Train on 1000 samples
1000/1000 [=====] - 0s 196us/sample - loss: 3.2446 - accuracy: 0.0940
Train on 1000 samples
1000/1000 [=====] - 0s 201us/sample - loss: 3.2777 - accuracy: 0.1080
Train on 1000 samples
1000/1000 [=====] - 0s 242us/sample - loss: 3.0962 - accuracy: 0.0820
Train on 1000 samples
1000/1000 [=====] - 0s 239us/sample - loss: 3.2774 - accuracy: 0.0750
Train on 1000 samples
1000/1000 [=====] - 0s 236us/sample - loss: 3.2792 - accuracy: 0.0900
Train on 1000 samples
1000/1000 [=====] - 0s 232us/sample - loss: 3.2315 - accuracy: 0.0980
Train on 1000 samples
1000/1000 [=====] - 0s 250us/sample - loss: 3.2029 - accuracy: 0.1020
Train on 1000 samples
1000/1000 [=====] - 0s 203us/sample - loss: 3.2055 - accuracy: 0.0930
Train on 1000 samples
1000/1000 [=====] - 0s 195us/sample - loss: 3.2551 - accuracy: 0.0840
Train on 1000 samples
1000/1000 [=====] - 0s 204us/sample - loss: 3.2564 - accuracy: 0.0930
Train on 1000 samples
1000/1000 [=====] - 0s 195us/sample - loss: 3.2150 - accuracy: 0.0910
```



```
In [45]: def neural_net_3hidden(l2):

    model = tf.keras.Sequential()
    if l2==True:
        model.add(tf.keras.layers.Dense(30, input_dim=784,activation
n='sigmoid',kernel_regularizer=tf.keras.regularizers.l2(5)))
        model.add(tf.keras.layers.Dense(30, input_dim=30,activation
='sigmoid',kernel_regularizer=tf.keras.regularizers.l2(5)))
        model.add(tf.keras.layers.Dense(30, input_dim=30,activation
='sigmoid',kernel_regularizer=tf.keras.regularizers.l2(5)))

    else:
        model.add(tf.keras.layers.Dense(30, input_dim=784,activation
n='sigmoid'))
        model.add(tf.keras.layers.Dense(30, input_dim=30,activation
='sigmoid'))
        model.add(tf.keras.layers.Dense(30, input_dim=30,activation
='sigmoid'))

    model.add(tf.keras.layers.Dense(10, activation='softmax'))

    model.summary()

    adam=tf.keras.optimizers.Adam(learning_rate=0.1)
    model.compile(loss='categorical_crossentropy', optimizer=adam,
metrics=[ 'accuracy' ])

    return model
```

3 hidden layers without regularization

Plot for Criterion Function Loss

```
In [46]: model=neural_net_3hidden(False)
history = model.fit(mini_train, train_labels, epochs=30, batch_size
=10,validation_data=(mini_test,test_labels))
model_training_plot()
```

Model: "sequential_15"

| Layer (type) | Output Shape | Param # |
|--------------------------|--------------|---------|
| dense_36 (Dense) | (None, 30) | 23550 |
| dense_37 (Dense) | (None, 30) | 930 |
| dense_38 (Dense) | (None, 30) | 930 |
| dense_39 (Dense) | (None, 10) | 310 |
| Total params: 25,720 | | |
| Trainable params: 25,720 | | |
| Non-trainable params: 0 | | |

Train on 1000 samples, validate on 1000 samples

Epoch 1/30

1000/1000 [=====] - 1s 883us/sample - loss: 2.3430 - accuracy: 0.1030 - val_loss: 2.0410 - val_accuracy: 0.2320

Epoch 2/30

1000/1000 [=====] - 0s 341us/sample - loss: 1.8662 - accuracy: 0.2940 - val_loss: 1.8509 - val_accuracy: 0.3050

Epoch 3/30

1000/1000 [=====] - 0s 337us/sample - loss: 1.5957 - accuracy: 0.3840 - val_loss: 1.5722 - val_accuracy: 0.3920

Epoch 4/30

1000/1000 [=====] - 0s 325us/sample - loss: 1.4944 - accuracy: 0.4430 - val_loss: 1.7053 - val_accuracy: 0.4120

Epoch 5/30

1000/1000 [=====] - 0s 334us/sample - loss: 1.3993 - accuracy: 0.4900 - val_loss: 1.6245 - val_accuracy: 0.4680

Epoch 6/30

1000/1000 [=====] - 0s 323us/sample - loss: 1.4007 - accuracy: 0.4990 - val_loss: 1.5509 - val_accuracy: 0.4680

Epoch 7/30

1000/1000 [=====] - 0s 322us/sample - loss: 1.3939 - accuracy: 0.4900 - val_loss: 1.5316 - val_accuracy: 0.4140


Epoch 8/30


1000/1000 [=====] - 0s 322us/sample - loss: 1.3420 - accuracy: 0.4700 - val_loss: 1.5455 - val_accuracy: 0.

```
4410
Epoch 9/30
1000/1000 [=====] - 0s 321us/sample - loss: 1.3711 - accuracy: 0.5030 - val_loss: 1.5507 - val_accuracy: 0.4520
Epoch 10/30
1000/1000 [=====] - 0s 328us/sample - loss: 1.2922 - accuracy: 0.5410 - val_loss: 1.4947 - val_accuracy: 0.4930
Epoch 11/30
1000/1000 [=====] - 0s 322us/sample - loss: 1.3013 - accuracy: 0.5040 - val_loss: 1.5241 - val_accuracy: 0.4650
Epoch 12/30
1000/1000 [=====] - 0s 333us/sample - loss: 1.2383 - accuracy: 0.5640 - val_loss: 1.4754 - val_accuracy: 0.4890
Epoch 13/30
1000/1000 [=====] - 0s 322us/sample - loss: 1.3024 - accuracy: 0.5410 - val_loss: 1.5744 - val_accuracy: 0.4690
Epoch 14/30
1000/1000 [=====] - 0s 327us/sample - loss: 1.2440 - accuracy: 0.5570 - val_loss: 1.4361 - val_accuracy: 0.5190
Epoch 15/30
1000/1000 [=====] - 0s 345us/sample - loss: 1.1231 - accuracy: 0.6300 - val_loss: 1.5081 - val_accuracy: 0.5040
Epoch 16/30
1000/1000 [=====] - 0s 328us/sample - loss: 1.0965 - accuracy: 0.6360 - val_loss: 1.5418 - val_accuracy: 0.4860
Epoch 17/30
1000/1000 [=====] - 0s 325us/sample - loss: 1.0873 - accuracy: 0.6460 - val_loss: 1.4526 - val_accuracy: 0.5360
Epoch 18/30
1000/1000 [=====] - 0s 328us/sample - loss: 1.0923 - accuracy: 0.6460 - val_loss: 1.4353 - val_accuracy: 0.5810
Epoch 19/30
1000/1000 [=====] - 0s 414us/sample - loss: 1.0920 - accuracy: 0.6490 - val_loss: 1.5757 - val_accuracy: 0.4940
Epoch 20/30
1000/1000 [=====] - 0s 378us/sample - loss: 1.1236 - accuracy: 0.6620 - val_loss: 1.3591 - val_accuracy: 0.5930
Epoch 21/30
1000/1000 [=====] - 0s 325us/sample - loss: 1.0635 - accuracy: 0.6630 - val_loss: 1.3930 - val_accuracy: 0.5770
```



```
Epoch 22/30
1000/1000 [=====] - 0s 347us/sample - loss: 1.1496 - accuracy: 0.6520 - val_loss: 1.3767 - val_accuracy: 0.5710
Epoch 23/30
1000/1000 [=====] - 0s 330us/sample - loss: 1.1327 - accuracy: 0.6360 - val_loss: 1.3660 - val_accuracy: 0.5840
Epoch 24/30
1000/1000 [=====] - 0s 453us/sample - loss: 1.0993 - accuracy: 0.6620 - val_loss: 1.4085 - val_accuracy: 0.5550
Epoch 25/30
1000/1000 [=====] - 0s 448us/sample - loss: 1.1731 - accuracy: 0.6380 - val_loss: 1.3314 - val_accuracy: 0.5870
Epoch 26/30
1000/1000 [=====] - 0s 460us/sample - loss: 1.0761 - accuracy: 0.6700 - val_loss: 1.3563 - val_accuracy: 0.5950
Epoch 27/30
1000/1000 [=====] - 0s 381us/sample - loss: 1.0148 - accuracy: 0.6760 - val_loss: 1.3234 - val_accuracy: 0.5920
Epoch 28/30
1000/1000 [=====] - 0s 362us/sample - loss: 1.0191 - accuracy: 0.6780 - val_loss: 1.3827 - val_accuracy: 0.5980
Epoch 29/30
1000/1000 [=====] - 0s 425us/sample - loss: 1.0864 - accuracy: 0.6470 - val_loss: 1.4749 - val_accuracy: 0.5400
Epoch 30/30
1000/1000 [=====] - 0s 372us/sample - loss: 1.0728 - accuracy: 0.6640 - val_loss: 1.4711 - val_accuracy: 0.5530
```





Plot for zero-one error

```
In [47]: model=neural_net_3hidden(False)
         traine,teste,traina,testa=for_error_plot(model)
         zero_one_error_plot(traine,teste,traina,testa)
```

Model: "sequential_16"

| Layer (type) | Output Shape | Param # |
|------------------|--------------|---------|
| dense_40 (Dense) | (None, 30) | 23550 |
| dense_41 (Dense) | (None, 30) | 930 |
| dense_42 (Dense) | (None, 30) | 930 |
| dense_43 (Dense) | (None, 10) | 310 |

Total params: 25,720

Trainable params: 25,720

Non-trainable params: 0

Train on 1000 samples

1000/1000 [=====] - 1s 631us/sample - loss: 2.1298 - accuracy: 0.1890

Train on 1000 samples

1000/1000 [=====] - 0s 201us/sample - loss: 1.5915 - accuracy: 0.3680

Train on 1000 samples

1000/1000 [=====] - 0s 204us/sample - loss: 1.4128 - accuracy: 0.4750

Train on 1000 samples


1000/1000 [=====] - 0s 261us/sample - loss: 1.3010 - accuracy: 0.5070


```
Train on 1000 samples
1000/1000 [=====] - 0s 222us/sample - loss: 1.2865 - accuracy: 0.5560
Train on 1000 samples
1000/1000 [=====] - 0s 199us/sample - loss: 1.2003 - accuracy: 0.5720
Train on 1000 samples
1000/1000 [=====] - 0s 211us/sample - loss: 1.1844 - accuracy: 0.5950
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 1.2122 - accuracy: 0.5820
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 1.2061 - accuracy: 0.5990
Train on 1000 samples
1000/1000 [=====] - 0s 200us/sample - loss: 1.2653 - accuracy: 0.5750
Train on 1000 samples
1000/1000 [=====] - 0s 198us/sample - loss: 1.2113 - accuracy: 0.5790
Train on 1000 samples
1000/1000 [=====] - 0s 199us/sample - loss: 1.1095 - accuracy: 0.6410
Train on 1000 samples
1000/1000 [=====] - 0s 370us/sample - loss: 1.0905 - accuracy: 0.6270
Train on 1000 samples
1000/1000 [=====] - 0s 198us/sample - loss: 1.0379 - accuracy: 0.6690
Train on 1000 samples
1000/1000 [=====] - 0s 207us/sample - loss: 1.0170 - accuracy: 0.6720
Train on 1000 samples
1000/1000 [=====] - 0s 199us/sample - loss: 1.0812 - accuracy: 0.6500
Train on 1000 samples
1000/1000 [=====] - 0s 208us/sample - loss: 1.0299 - accuracy: 0.6700
Train on 1000 samples
1000/1000 [=====] - 0s 221us/sample - loss: 1.0339 - accuracy: 0.6890
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 1.0752 - accuracy: 0.6800
Train on 1000 samples
1000/1000 [=====] - 0s 282us/sample - loss: 1.0100 - accuracy: 0.7100
Train on 1000 samples
1000/1000 [=====] - 0s 205us/sample - loss: 0.9543 - accuracy: 0.7100
Train on 1000 samples
1000/1000 [=====] - 0s 220us/sample - loss:
```

```

s: 0.9570 - accuracy: 0.6930
Train on 1000 samples
1000/1000 [=====] - 0s 221us/sample - los
s: 0.9696 - accuracy: 0.7160
Train on 1000 samples
1000/1000 [=====] - 0s 203us/sample - los
s: 0.9250 - accuracy: 0.7280
Train on 1000 samples
1000/1000 [=====] - 0s 251us/sample - los
s: 0.9687 - accuracy: 0.7160
Train on 1000 samples
1000/1000 [=====] - 0s 206us/sample - los
s: 1.0204 - accuracy: 0.6930
Train on 1000 samples
1000/1000 [=====] - 0s 229us/sample - los
s: 0.9849 - accuracy: 0.7180
Train on 1000 samples
1000/1000 [=====] - 0s 198us/sample - los
s: 0.9518 - accuracy: 0.7130
Train on 1000 samples
1000/1000 [=====] - 0s 205us/sample - los
s: 1.0698 - accuracy: 0.6690
Train on 1000 samples
1000/1000 [=====] - 0s 229us/sample - los
s: 0.9762 - accuracy: 0.7100

```





Plot for learning rate

```
In [48]: model=neural_net_3hidden(False)
         learning_rate_plot(25720,model)
```

Model: "sequential_17"

| Layer (type) | Output Shape | Param # |
|--------------------------|--------------|---------|
| dense_44 (Dense) | (None, 30) | 23550 |
| dense_45 (Dense) | (None, 30) | 930 |
| dense_46 (Dense) | (None, 30) | 930 |
| dense_47 (Dense) | (None, 10) | 310 |
| Total params: 25,720 | | |
| Trainable params: 25,720 | | |
| Non-trainable params: 0 | | |

Train on 1000 samples

1000/1000 [=====] - 1s 647us/sample - loss: 2.1075 - accuracy: 0.2030

Train on 1000 samples

1000/1000 [=====] - 0s 228us/sample - loss: 1.5866 - accuracy: 0.3830

Train on 1000 samples

1000/1000 [=====] - 0s 200us/sample - loss: 1.4277 - accuracy: 0.4840

Train on 1000 samples

1000/1000 [=====] - 0s 210us/sample - loss: 1.3228 - accuracy: 0.5160


Train on 1000 samples

```
1000/1000 [=====] - 0s 198us/sample - loss: 1.2775 - accuracy: 0.5270
Train on 1000 samples
1000/1000 [=====] - 0s 226us/sample - loss: 1.3131 - accuracy: 0.5260
Train on 1000 samples
1000/1000 [=====] - 0s 250us/sample - loss: 1.2109 - accuracy: 0.5840
Train on 1000 samples
1000/1000 [=====] - 0s 248us/sample - loss: 1.1113 - accuracy: 0.6450
Train on 1000 samples
1000/1000 [=====] - 0s 278us/sample - loss: 1.1134 - accuracy: 0.6090
Train on 1000 samples
1000/1000 [=====] - 0s 246us/sample - loss: 1.2456 - accuracy: 0.5790
Train on 1000 samples
1000/1000 [=====] - 0s 263us/sample - loss: 1.1719 - accuracy: 0.5980
Train on 1000 samples
1000/1000 [=====] - 0s 262us/sample - loss: 1.1489 - accuracy: 0.6250
Train on 1000 samples
1000/1000 [=====] - 0s 266us/sample - loss: 1.1796 - accuracy: 0.6010
Train on 1000 samples
1000/1000 [=====] - 0s 223us/sample - loss: 1.1008 - accuracy: 0.6580
Train on 1000 samples
1000/1000 [=====] - 0s 258us/sample - loss: 1.1574 - accuracy: 0.6280
Train on 1000 samples
1000/1000 [=====] - 0s 256us/sample - loss: 1.1407 - accuracy: 0.6300
Train on 1000 samples
1000/1000 [=====] - 0s 240us/sample - loss: 1.0512 - accuracy: 0.6790
Train on 1000 samples
1000/1000 [=====] - 0s 237us/sample - loss: 0.9877 - accuracy: 0.6960
Train on 1000 samples
1000/1000 [=====] - 0s 204us/sample - loss: 0.9734 - accuracy: 0.6920
Train on 1000 samples
1000/1000 [=====] - 0s 235us/sample - loss: 1.0383 - accuracy: 0.6700
Train on 1000 samples
1000/1000 [=====] - 0s 203us/sample - loss: 1.0989 - accuracy: 0.6620
Train on 1000 samples
1000/1000 [=====] - 0s 200us/sample - loss: 1.1810 - accuracy: 0.6100
```

```

Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 1.0507 - accuracy: 0.6680
Train on 1000 samples
1000/1000 [=====] - 0s 212us/sample - loss: 1.0581 - accuracy: 0.6760
Train on 1000 samples
1000/1000 [=====] - 0s 257us/sample - loss: 1.0162 - accuracy: 0.6810
Train on 1000 samples
1000/1000 [=====] - 0s 248us/sample - loss: 0.9451 - accuracy: 0.7000
Train on 1000 samples
1000/1000 [=====] - 0s 245us/sample - loss: 1.0069 - accuracy: 0.6870
Train on 1000 samples
1000/1000 [=====] - 0s 246us/sample - loss: 0.9982 - accuracy: 0.6950
Train on 1000 samples
1000/1000 [=====] - 0s 263us/sample - loss: 1.0097 - accuracy: 0.6760
Train on 1000 samples
1000/1000 [=====] - 0s 235us/sample - loss: 1.0494 - accuracy: 0.6540

```



3 hidden layers with regularization

```

In [49]: model=neural_net_3hidden(True)
history = model.fit(mini_train, train_labels, epochs=30, batch_size=10, validation_data=(mini_test, test_labels))

```

Model: "sequential_18"

| Layer (type) | Output Shape | Param # |
|--------------|--------------|---------|
| ===== | | |

| | | |
|------------------|------------|-------|
| dense_48 (Dense) | (None, 30) | 23550 |
| dense_49 (Dense) | (None, 30) | 930 |
| dense_50 (Dense) | (None, 30) | 930 |
| dense_51 (Dense) | (None, 10) | 310 |

=====
Total params: 25,720

Trainable params: 25,720

Non-trainable params: 0

=====
Train on 1000 samples, validate on 1000 samples

Epoch 1/30

1000/1000 [=====] - 1s 1ms/sample - loss: 40.2314 - accuracy: 0.0870 - val_loss: 2.3294 - val_accuracy: 0.1000

Epoch 2/30

1000/1000 [=====] - 0s 335us/sample - loss: 2.3410 - accuracy: 0.0870 - val_loss: 2.3274 - val_accuracy: 0.1000

Epoch 3/30

1000/1000 [=====] - 0s 339us/sample - loss: 2.3343 - accuracy: 0.0880 - val_loss: 2.3309 - val_accuracy: 0.1000

Epoch 4/30

1000/1000 [=====] - 0s 337us/sample - loss: 2.3348 - accuracy: 0.0950 - val_loss: 2.3154 - val_accuracy: 0.1000

Epoch 5/30

1000/1000 [=====] - 0s 331us/sample - loss: 2.3340 - accuracy: 0.1010 - val_loss: 2.3240 - val_accuracy: 0.1000

Epoch 6/30

1000/1000 [=====] - 0s 326us/sample - loss: 2.3341 - accuracy: 0.0870 - val_loss: 2.3132 - val_accuracy: 0.1000

Epoch 7/30

1000/1000 [=====] - 0s 326us/sample - loss: 2.3319 - accuracy: 0.0770 - val_loss: 2.3103 - val_accuracy: 0.1000

Epoch 8/30

1000/1000 [=====] - 0s 324us/sample - loss: 2.3252 - accuracy: 0.0900 - val_loss: 2.3129 - val_accuracy: 0.1000

Epoch 9/30

1000/1000 [=====] - 0s 330us/sample - loss: 2.3321 - accuracy: 0.0850 - val_loss: 2.3105 - val_accuracy: 0.1000

Epoch 10/30



1000/1000 [=====] - 0s 349us/sample - loss: 2.3332 - accuracy: 0.1040 - val_loss: 2.3180 - val_accuracy: 0.1000


```
Epoch 11/30
1000/1000 [=====] - 0s 353us/sample - loss: 2.3272 - accuracy: 0.0830 - val_loss: 2.3172 - val_accuracy: 0.1000
Epoch 12/30
1000/1000 [=====] - 0s 341us/sample - loss: 2.3325 - accuracy: 0.0890 - val_loss: 2.3182 - val_accuracy: 0.1000
Epoch 13/30
1000/1000 [=====] - 0s 442us/sample - loss: 3.1256 - accuracy: 0.1020 - val_loss: 3.8704 - val_accuracy: 0.1000
Epoch 14/30
1000/1000 [=====] - 0s 434us/sample - loss: 3.3208 - accuracy: 0.0830 - val_loss: 2.9531 - val_accuracy: 0.1000
Epoch 15/30
1000/1000 [=====] - 0s 418us/sample - loss: 2.9627 - accuracy: 0.1070 - val_loss: 2.9554 - val_accuracy: 0.1000
Epoch 16/30
1000/1000 [=====] - 0s 426us/sample - loss: 2.9534 - accuracy: 0.0970 - val_loss: 2.9018 - val_accuracy: 0.1000
Epoch 17/30
1000/1000 [=====] - 0s 432us/sample - loss: 2.9412 - accuracy: 0.1010 - val_loss: 3.0722 - val_accuracy: 0.1000
Epoch 18/30
1000/1000 [=====] - 0s 333us/sample - loss: 3.4950 - accuracy: 0.0920 - val_loss: 3.9782 - val_accuracy: 0.1000
Epoch 19/30
1000/1000 [=====] - 0s 344us/sample - loss: 3.5994 - accuracy: 0.0940 - val_loss: 3.3505 - val_accuracy: 0.1000
Epoch 20/30
1000/1000 [=====] - 0s 343us/sample - loss: 3.1450 - accuracy: 0.1080 - val_loss: 3.0348 - val_accuracy: 0.1000
Epoch 21/30
1000/1000 [=====] - 0s 329us/sample - loss: 3.0183 - accuracy: 0.0950 - val_loss: 2.9622 - val_accuracy: 0.1000
Epoch 22/30
1000/1000 [=====] - 0s 332us/sample - loss: 3.2708 - accuracy: 0.0960 - val_loss: 3.4565 - val_accuracy: 0.1000
Epoch 23/30
1000/1000 [=====] - 0s 438us/sample - loss: 3.4836 - accuracy: 0.0900 - val_loss: 3.3276 - val_accuracy: 0.1000
Epoch 24/30
```

```
1000/1000 [=====] - 0s 378us/sample - loss: 3.2950 - accuracy: 0.0990 - val_loss: 3.1120 - val_accuracy: 0.1000
Epoch 25/30
1000/1000 [=====] - 0s 330us/sample - loss: 3.1229 - accuracy: 0.0880 - val_loss: 3.2092 - val_accuracy: 0.1000
Epoch 26/30
1000/1000 [=====] - 0s 344us/sample - loss: 3.3273 - accuracy: 0.0950 - val_loss: 3.4942 - val_accuracy: 0.1000
Epoch 27/30
1000/1000 [=====] - 0s 337us/sample - loss: 3.3847 - accuracy: 0.0890 - val_loss: 3.2494 - val_accuracy: 0.1000
Epoch 28/30
1000/1000 [=====] - 0s 327us/sample - loss: 3.2190 - accuracy: 0.0930 - val_loss: 3.1648 - val_accuracy: 0.1000
Epoch 29/30
1000/1000 [=====] - 0s 337us/sample - loss: 3.2281 - accuracy: 0.0940 - val_loss: 3.2706 - val_accuracy: 0.1000
Epoch 30/30
1000/1000 [=====] - 0s 329us/sample - loss: 3.3710 - accuracy: 0.0980 - val_loss: 3.4740 - val_accuracy: 0.1000
```

Plot for criterion function loss

```
In [50]: model_training_plot()
```



Plot for zero-one error

```
In [51]: model=neural_net_3hidden(True)
         traine,teste,traina,testa=for_error_plot(model)
         zero_one_error_plot(traine,teste,traina,testa)
```


Model: "sequential_19"

| Layer (type) | Output Shape | Param # |
|------------------|--------------|---------|
| dense_52 (Dense) | (None, 30) | 23550 |
| dense_53 (Dense) | (None, 30) | 930 |
| dense_54 (Dense) | (None, 30) | 930 |
| dense_55 (Dense) | (None, 10) | 310 |

```
=====
Total params: 25,720
Trainable params: 25,720
Non-trainable params: 0

Train on 1000 samples
1000/1000 [=====] - 1s 750us/sample - loss: 40.4962 - accuracy: 0.1090
Train on 1000 samples
1000/1000 [=====] - 0s 226us/sample - loss: 2.3453 - accuracy: 0.1090
Train on 1000 samples
1000/1000 [=====] - 0s 280us/sample - loss: 2.3412 - accuracy: 0.0980
Train on 1000 samples
1000/1000 [=====] - 0s 257us/sample - loss: 2.3324 - accuracy: 0.1150
Train on 1000 samples
1000/1000 [=====] - 0s 198us/sample - loss: 2.3385 - accuracy: 0.0850
Train on 1000 samples
1000/1000 [=====] - 0s 272us/sample - loss: 2.3301 - accuracy: 0.0960
Train on 1000 samples
1000/1000 [=====] - 0s 246us/sample - loss: 2.3243 - accuracy: 0.0950
Train on 1000 samples
1000/1000 [=====] - 0s 212us/sample - loss: 2.3283 - accuracy: 0.1030
Train on 1000 samples
1000/1000 [=====] - 0s 246us/sample - loss: 2.3264 - accuracy: 0.0980
Train on 1000 samples
1000/1000 [=====] - 0s 233us/sample - loss: 2.3368 - accuracy: 0.0950
Train on 1000 samples
1000/1000 [=====] - 0s 204us/sample - loss: 2.3364 - accuracy: 0.1090
Train on 1000 samples
1000/1000 [=====] - 0s 207us/sample - loss: 2.3298 - accuracy: 0.0780
Train on 1000 samples
1000/1000 [=====] - 0s 203us/sample - loss: 3.1083 - accuracy: 0.0900
Train on 1000 samples
1000/1000 [=====] - 0s 207us/sample - loss: 3.3003 - accuracy: 0.0820
Train on 1000 samples
1000/1000 [=====] - 0s 204us/sample - loss: 2.9490 - accuracy: 0.0970
Train on 1000 samples
1000/1000 [=====] - 0s 229us/sample - loss: 2.9410 - accuracy: 0.0870
```

```
Train on 1000 samples
1000/1000 [=====] - 0s 203us/sample - loss: 2.9352 - accuracy: 0.0870
Train on 1000 samples
1000/1000 [=====] - 0s 207us/sample - loss: 3.5555 - accuracy: 0.0890
Train on 1000 samples
1000/1000 [=====] - 0s 206us/sample - loss: 3.5269 - accuracy: 0.0780
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 3.1053 - accuracy: 0.1070
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 3.0694 - accuracy: 0.0840
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 3.3600 - accuracy: 0.0900
Train on 1000 samples
1000/1000 [=====] - 0s 203us/sample - loss: 3.4318 - accuracy: 0.0970
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 3.2234 - accuracy: 0.0910
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 3.1541 - accuracy: 0.0940
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 3.4287 - accuracy: 0.0960
Train on 1000 samples
1000/1000 [=====] - 0s 207us/sample - loss: 3.3139 - accuracy: 0.0960
Train on 1000 samples
1000/1000 [=====] - 0s 204us/sample - loss: 3.1531 - accuracy: 0.0940
Train on 1000 samples
1000/1000 [=====] - 0s 206us/sample - loss: 3.2569 - accuracy: 0.1010
Train on 1000 samples
1000/1000 [=====] - 0s 203us/sample - loss: 3.4481 - accuracy: 0.0880
```



Plot for learning rate

```
In [52]: model=neural_net_3hidden(True)
         learning_rate_plot(25720,model)
```


Model: "sequential_20"

| Layer (type) | Output Shape | Param # |
|------------------|--------------|---------|
| dense_56 (Dense) | (None, 30) | 23550 |
| dense_57 (Dense) | (None, 30) | 930 |
| dense_58 (Dense) | (None, 30) | 930 |
| dense_59 (Dense) | (None, 10) | 310 |

Total params: 25,720
Trainable params: 25,720
Non-trainable params: 0

Train on 1000 samples
1000/1000 [=====] - 1s 743us/sample - loss: 40.6000 - accuracy: 0.0870
Train on 1000 samples
1000/1000 [=====] - 0s 213us/sample - loss: 2.3403 - accuracy: 0.0950
Train on 1000 samples
1000/1000 [=====] - 0s 253us/sample - loss: 2.3368 - accuracy: 0.0900s - loss: 2.3385 - accuracy: 0.
Train on 1000 samples
1000/1000 [=====] - 0s 248us/sample - loss: 2.3359 - accuracy: 0.0980
Train on 1000 samples
1000/1000 [=====] - 0s 223us/sample - loss: 2.3272 - accuracy: 0.0920
Train on 1000 samples
1000/1000 [=====] - 0s 201us/sample - loss: 2.3344 - accuracy: 0.0980
Train on 1000 samples
1000/1000 [=====] - 0s 219us/sample - loss: 2.3256 - accuracy: 0.0910
Train on 1000 samples
1000/1000 [=====] - 0s 212us/sample - loss: 2.3293 - accuracy: 0.0920
Train on 1000 samples
1000/1000 [=====] - 0s 203us/sample - loss: 2.3290 - accuracy: 0.0880
Train on 1000 samples
1000/1000 [=====] - 0s 291us/sample - loss: 2.3378 - accuracy: 0.0920
Train on 1000 samples
1000/1000 [=====] - 0s 209us/sample - loss: 2.3372 - accuracy: 0.0870
Train on 1000 samples
1000/1000 [=====] - 0s 281us/sample - loss: 2.3270 - accuracy: 0.0900
Train on 1000 samples
1000/1000 [=====] - 0s 206us/sample - loss: 3.1312 - accuracy: 0.0720
Train on 1000 samples
1000/1000 [=====] - 0s 204us/sample - loss: 3.2935 - accuracy: 0.0900
Train on 1000 samples
1000/1000 [=====] - 0s 221us/sample - loss: 2.9718 - accuracy: 0.0920
Train on 1000 samples
1000/1000 [=====] - 0s 201us/sample - loss: 2.9480 - accuracy: 0.1140
Train on 1000 samples

```
1000/1000 [=====] - 0s 203us/sample - loss: 2.9606 - accuracy: 0.0940
Train on 1000 samples
1000/1000 [=====] - 0s 201us/sample - loss: 3.5531 - accuracy: 0.1040
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 3.5676 - accuracy: 0.0730
Train on 1000 samples
1000/1000 [=====] - 0s 203us/sample - loss: 3.1163 - accuracy: 0.0990
Train on 1000 samples
1000/1000 [=====] - 0s 221us/sample - loss: 3.0721 - accuracy: 0.0950
Train on 1000 samples
1000/1000 [=====] - 0s 246us/sample - loss: 3.2771 - accuracy: 0.0930
Train on 1000 samples
1000/1000 [=====] - 0s 258us/sample - loss: 3.4924 - accuracy: 0.0840
Train on 1000 samples
1000/1000 [=====] - 0s 201us/sample - loss: 3.2208 - accuracy: 0.0980
Train on 1000 samples
1000/1000 [=====] - 0s 209us/sample - loss: 3.1684 - accuracy: 0.0940
Train on 1000 samples
1000/1000 [=====] - 0s 212us/sample - loss: 3.3135 - accuracy: 0.0940
Train on 1000 samples
1000/1000 [=====] - 0s 201us/sample - loss: 3.4041 - accuracy: 0.0920
Train on 1000 samples
1000/1000 [=====] - 0s 202us/sample - loss: 3.2010 - accuracy: 0.0780
Train on 1000 samples
1000/1000 [=====] - 0s 200us/sample - loss: 3.2152 - accuracy: 0.0880
Train on 1000 samples
1000/1000 [=====] - 0s 201us/sample - loss: 3.4112 - accuracy: 0.0860
```

Question 2)c) CNN

```
In [53]: from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Activation, Flatten
from keras.optimizers import Adam
from keras.layers.normalization import BatchNormalization
from keras.utils import np_utils
from keras.layers import Conv2D, MaxPooling2D, ZeroPadding2D, GlobalAveragePooling2D
from keras.layers.advanced_activations import LeakyReLU
from keras.preprocessing.image import ImageDataGenerator
```

```
In [64]: model = Sequential()

model.add(Conv2D(32, (3, 3), input_shape=(28,28,1)))
model.add(LeakyReLU(alpha=0.1))
model.add(Conv2D(32, (3, 3)))
model.add(LeakyReLU(alpha=0.1))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Conv2D(64, (3, 3)))
model.add(LeakyReLU(alpha=0.1))
model.add(Conv2D(64, (3, 3)))
model.add(LeakyReLU(alpha=0.1))
model.add(MaxPooling2D(pool_size=(2,2)))

model.add(Flatten())

model.add(Dense(512))
model.add(LeakyReLU(alpha=0.1))
model.add(Dropout(0.2))
model.add(Dense(10))

model.add(Activation('softmax'))
```

```
In [65]: model.summary()
```

Model: "sequential_2"

| Layer (type) | Output Shape | Param # |
|--------------------------------|--------------------|---------|
| conv2d_5 (Conv2D) | (None, 26, 26, 32) | 320 |
| leaky_re_lu_1 (LeakyReLU) | (None, 26, 26, 32) | 0 |
| conv2d_6 (Conv2D) | (None, 24, 24, 32) | 9248 |
| leaky_re_lu_2 (LeakyReLU) | (None, 24, 24, 32) | 0 |
| max_pooling2d_3 (MaxPooling2D) | (None, 12, 12, 32) | 0 |
| conv2d_7 (Conv2D) | (None, 10, 10, 64) | 18496 |
| leaky_re_lu_3 (LeakyReLU) | (None, 10, 10, 64) | 0 |
| conv2d_8 (Conv2D) | (None, 8, 8, 64) | 36928 |
| leaky_re_lu_4 (LeakyReLU) | (None, 8, 8, 64) | 0 |
| max_pooling2d_4 (MaxPooling2D) | (None, 4, 4, 64) | 0 |
| flatten_2 (Flatten) | (None, 1024) | 0 |
| dense_3 (Dense) | (None, 512) | 524800 |
| leaky_re_lu_5 (LeakyReLU) | (None, 512) | 0 |
| dropout_2 (Dropout) | (None, 512) | 0 |
| dense_4 (Dense) | (None, 10) | 5130 |
| activation_7 (Activation) | (None, 10) | 0 |
| Total params: 594,922 | | |
| Trainable params: 594,922 | | |
| Non-trainable params: 0 | | |

```
In [66]: mini_train=mini_train.reshape((1000,28,28,1))
mini_test=mini_test.reshape((1000,28,28,1))
```

```
In [67]: mini_train.shape,mini_test.shape
```

```
Out[67]: ((1000, 28, 28, 1), (1000, 28, 28, 1))
```

```
In [68]: train_labels.shape,test_labels.shape
```

```
Out[68]: ((1000, 10), (1000, 10))
```

```

In [69]: model.compile(loss='categorical_crossentropy', optimizer=Adam(learn
         ing_rate=0.1), metrics=['accuracy'])

In [70]: train_gen = ImageDataGenerator(rotation_range=3, width_shift_range=
         3, height_shift_range=3)

         test_gen = ImageDataGenerator()

In [71]: train_generator = train_gen.flow(mini_train, train_labels, batch_si
         ze=10)
         test_generator = test_gen.flow(mini_test, test_labels, batch_size=1
         0)

In [72]: history=model.fit_generator(train_generator, steps_per_epoch=1000//
         10, epochs=30, validation_data=test_generator, validation_steps=100
         0//10)


Epoch 1/30
100/100 [=====] - 5s 49ms/step - loss: 13
943718.7768 - accuracy: 0.1110 - val_loss: 16028184.0000 - val_acc
uracy: 0.1260
Epoch 2/30
100/100 [=====] - 4s 44ms/step - loss: 55
29694.5050 - accuracy: 0.2920 - val_loss: 788680.6250 - val_accura
cy: 0.3530
Epoch 3/30
100/100 [=====] - 4s 44ms/step - loss: 45
1844.4041 - accuracy: 0.5000 - val_loss: 159742.6250 - val_accurac
y: 0.6720
Epoch 4/30
100/100 [=====] - 4s 44ms/step - loss: 22
2572.9458 - accuracy: 0.5340 - val_loss: 352928.4062 - val_accurac
y: 0.7220
Epoch 5/30
100/100 [=====] - 4s 43ms/step - loss: 24
8917.9212 - accuracy: 0.5380 - val_loss: 259498.5000 - val_accurac
y: 0.5210
Epoch 6/30
100/100 [=====] - 5s 45ms/step - loss: 18
4655.6699 - accuracy: 0.5590 - val_loss: 90099.7891 - val_accuracy
: 0.6240
Epoch 7/30
100/100 [=====] - 5s 46ms/step - loss: 14
4813.8153 - accuracy: 0.6100 - val_loss: 25548.2871 - val_accuracy
: 0.7240
Epoch 8/30
100/100 [=====] - 5s 46ms/step - loss: 92
076.4708 - accuracy: 0.6470 - val_loss: 118837.2266 - val_accuracy
: 0.7590
Epoch 9/30
100/100 [=====] - 5s 46ms/step - loss: 60

```

```
656.4339 - accuracy: 0.7140 - val_loss: 87568.8125 - val_accuracy:
0.7610
Epoch 10/30
100/100 [=====] - 5s 46ms/step - loss: 39
367.4359 - accuracy: 0.7390 - val_loss: 13037.2109 - val_accuracy:
0.7950
Epoch 11/30
100/100 [=====] - 5s 48ms/step - loss: 30
173.6685 - accuracy: 0.7420 - val_loss: 27648.9473 - val_accuracy:
0.7780
Epoch 12/30
100/100 [=====] - 5s 46ms/step - loss: 24
554.4343 - accuracy: 0.7470 - val_loss: 21893.1211 - val_accuracy:
0.8160
Epoch 13/30
100/100 [=====] - 5s 45ms/step - loss: 31
150.1836 - accuracy: 0.7590 - val_loss: 5697.6123 - val_accuracy:
0.7700
Epoch 14/30
100/100 [=====] - 5s 45ms/step - loss: 22
701.2793 - accuracy: 0.7860 - val_loss: 11643.3145 - val_accuracy:
0.8480
Epoch 15/30
100/100 [=====] - 5s 47ms/step - loss: 20
547.3794 - accuracy: 0.7990 - val_loss: 110618.9766 - val_accuracy
: 0.7430
Epoch 16/30
100/100 [=====] - 5s 47ms/step - loss: 20
918.9665 - accuracy: 0.8110 - val_loss: 34346.0859 - val_accuracy:
0.8690
Epoch 17/30
100/100 [=====] - 5s 46ms/step - loss: 20
114.8533 - accuracy: 0.8040 - val_loss: 3444.9390 - val_accuracy:
0.7690
Epoch 18/30
100/100 [=====] - 5s 47ms/step - loss: 15
939.2175 - accuracy: 0.8310 - val_loss: 0.0000e+00 - val_accuracy:
0.8720
Epoch 19/30
100/100 [=====] - 4s 44ms/step - loss: 28
562.0061 - accuracy: 0.7770 - val_loss: 52312.0234 - val_accuracy:
0.7820
Epoch 20/30
100/100 [=====] - 4s 43ms/step - loss: 23
577.5042 - accuracy: 0.7930 - val_loss: 24177.2129 - val_accuracy:
0.8570
Epoch 21/30
100/100 [=====] - 5s 46ms/step - loss: 14
180.6516 - accuracy: 0.8050 - val_loss: 9576.2246 - val_accuracy:
0.8380
Epoch 22/30
100/100 [=====] - 4s 45ms/step - loss: 21
023.0222 - accuracy: 0.7800 - val_loss: 51836.2148 - val_accuracy:
```

```
0.8530
Epoch 23/30
100/100 [=====] - 5s 45ms/step - loss: 29
084.9780 - accuracy: 0.7730 - val_loss: 56501.9062 - val_accuracy:
0.8600
Epoch 24/30
100/100 [=====] - 5s 45ms/step - loss: 21
363.0631 - accuracy: 0.8190 - val_loss: 11489.7939 - val_accuracy:
0.8370
Epoch 25/30
100/100 [=====] - 5s 45ms/step - loss: 14
810.4845 - accuracy: 0.8210 - val_loss: 0.0000e+00 - val_accuracy:
0.8630
Epoch 26/30
100/100 [=====] - 5s 51ms/step - loss: 16
678.5440 - accuracy: 0.8170 - val_loss: 15532.1309 - val_accuracy:
0.8450
Epoch 27/30
100/100 [=====] - 5s 47ms/step - loss: 15
543.7392 - accuracy: 0.8240 - val_loss: 0.0000e+00 - val_accuracy:
0.8750
Epoch 28/30
100/100 [=====] - 5s 48ms/step - loss: 17
396.9999 - accuracy: 0.8280 - val_loss: 29576.3809 - val_accuracy:
0.7420
Epoch 29/30
100/100 [=====] - 4s 42ms/step - loss: 25
180.8621 - accuracy: 0.7610 - val_loss: 0.0000e+00 - val_accuracy:
0.8310
Epoch 30/30
100/100 [=====] - 4s 43ms/step - loss: 49
357.7607 - accuracy: 0.7530 - val_loss: 120708.9219 - val_accuracy
: 0.7700
```

```
In [73]: model_training_plot()
```



References :

<https://yashk2810.github.io/Applying-Convolutional-Neural-Network-on-the-MNIST-dataset/>

(<https://yashk2810.github.io/Applying-Convolutional-Neural-Network-on-the-MNIST-dataset/>)

<https://rohanvarma.me/Regularization/> (<https://rohanvarma.me/Regularization/>)

<https://ljvmiranda921.github.io/notebook/2017/08/13/softmax-and-the-negative-log-likelihood/>

(<https://ljvmiranda921.github.io/notebook/2017/08/13/softmax-and-the-negative-log-likelihood/>)